

**PRACTICAL JOURNAL IN
ADVANCED ARTIFICIAL INTELLIGENCE
MACHINE LEARNING**

SUBMITTED BY

SONAWANE CHAITANYA RAJESH

ROLL NO: 2414559

**IN PARTIAL FULLFILMENT FOR THE DEGREE OF
MASTER OF SCIENCE IN INFORMATION TECHNOLOGY PART – II
SEMESTER III**

ACADEMIC YEAR

2024-2025

॥ आ नो भद्राः क्रतवो यन्तु विश्वतः ॥



**PARLE TILAK VIDYALAYA ASSOCIATION'S
MULUND COLLEGE OF COMMERCE(AUTONOMOUS)
(AFFILIATED TO UNIVERSITY OF MUMBAI)
NAAC RE-ACCREDITED A GRADE – III CYCLE
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA
2024-25**



Parle Tilak Vidyalaya Association's
MULUND COLLEGE OF COMMERCE
(AUTONOMOUS)

(Affiliated to University of Mumbai)
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that SONAWANE CHAITANYA RAJESH of **M.Sc. I.T. Part II** Roll No 2414559 has successfully completed the practical work in Advanced Artificial Intelligence in partial fulfilment of the requirements for the Semester III of **M.Sc. I.T. Part II** during the academic year **2024-25**.

Teacher In-charge and Coordinator

Examiner

Date:

College Seal

INDEX

Practical No.	Date	Name of Practical	Page No.	Signature
1	06/10/2022	Design a bot using AIML	3	
2	13/10/2022	Design an Expert system using AIML	7	
3	20/10/2022	Implement Bayes Theorem using Python	11	
4	27/10/2022	Implement Conditional Probability and joint probability using Python.	12	
5	03/11/2022	Write a program for to implement Rule based system. (Prolog)	14	
6	10/11/2022	Design a Fuzzy based application using Python / R	18	
7	17/11/2022	[A] Write an application to stimulate supervised learning model.	24	
	24/11/2022	[B] Write an application to stimulate unsupervised learning model.	26	
8	01/12/2022	Write an application to implement clustering algorithm.	28	
9	08/12/2022	Write an application to implement support vector machine algorithm.	30	
10	15/12/2022	Simulate artificial neural network model with both feedforward and backpropagation approach.	34	

PRACTICAL 1

AIM: Design a bot using AIML.

CODE:

Step 1: Create the XML file.

Open the notepad, write the following code, and save it as std-startup.xml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- std-startup.xml -->
  <!-- Category is an atomic AIML unit -->
  <category>
    <!-- Pattern to match in user input -->
    <!-- If user enters "LOAD AIML B" -->
    <pattern>LOAD AIML B</pattern>
    <!-- Template is the response to the pattern -->
    <!-- This learn an aiml file -->
    <template>
      <learn>basic_chat.aiml</learn>
      <!-- You can add more aiml files here -->
      <!--<learn>more_aiml.aiml</learn>-->
    </template>
  </category>
</aiml>
```

Step 2: Create the aiml file.

Open the notepad, write the following code, and save it as basic_chat.aiml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- basic_chat.aiml -->
  <category>
    <pattern>HELLO</pattern>
    <template>
      Well, hello!
    </template>
  </category>
  <category>
    <pattern>WHAT ARE YOU</pattern>
```

```
<template> I'm a bot, silly! </template>
</category>
<category>
  <pattern>MY NAME IS *</pattern>
  <template>
    <set name = "username">
      <star/>
    </set> is the nice name.
  </template>
</category>
<category>
  <pattern>I LIKE *</pattern>
  <template>
    <set name = "liking">
      <star/>
    </set> is also my favourite.
  </template>
</category>
<category>
  <pattern>MY DOG NAME IS *</pattern>
  <template>
    THAT IS INTERESTING THAT YOU HAVE A DOG NAMED
      <set name ="dog">
        <star/>
      </set> .
    </template>
  </category>
  <category>
    <pattern>BYE</pattern>
    <template>
      Bye!!!
      <get name = "username"/> Thanks for talking with me.
    </template>
  </category>
</aiml>
```

Step 3: Install aiml packages

```
pip install aiml
pip install aimlbotkernel
or
pip3 install aiml
pip3 install aimlbotkernel
```

Step 4: Create chatbot.py file

```
import aiml # Create the kernel and learn AIML files
kernel = aiml.Kernel()
kernel.learn("std-startup.xml")
kernel.respond("load aiml b") # Press CTRL-C to break this loop
while True:
    message = input("Enter your message to the bot: ")
    if message == "quit":
        break
    else:
        bot_response = kernel.respond(message)
        print(bot_response)
```

OUTPUT:

```
Loading std-startup.xml...done (0.03 seconds)
Loading basic_chat.aiml...done (0.00 seconds)
Enter your message to the bot: Hello
Well, hello!
Enter your message to the bot: What are you
I'm a bot, silly!
Enter your message to the bot: My name is Prateek
Prateek is the nice name.
Enter your message to the bot: I like AIML
AIML is also my favourite.
Enter your message to the bot: My dog name is Rex
THAT IS INTERESTING THAT YOU HAVE A DOG NAMED Rex
Enter your message to the bot: Bye
Bye!!! Prateek Thanks for talking with me.
Enter your message to the bot:
```

PRACTICAL 2

AIM: Design an Expert system using AIML.

CODE:

Step 1: Create the XML file

Open the notepad, write the following code, and save it as std-startup.xml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- std-startup.xml -->
  <!-- Category is an atomic AIML unit -->
  <category>
    <!-- Pattern to match in user input -->
    <!-- If user enters "LOAD AIML B" -->
    <pattern>LOAD AIML B</pattern>
    <!-- Template is the response to the pattern -->
    <!-- This learn an aiml file -->
    <template>
      <learn>basic_chat.aiml</learn>
      <!-- You can add more aiml files here -->
      <!--<learn>more_aiml.aiml</learn>-->
    </template>
  </category>
</aiml>
```

Step 2: Create the aiml file

Open the notepad, write the following code, and save it as basic_chat.aiml

```
<aiml version="1.0.1" encoding="UTF-8">
<!-- basic_chat.aiml -->
  <category>
    <pattern>HELLO</pattern>
    <template>
      WHAT WOULD YOU LIKE TO DISCUSS? : HEALTH, MOVIES
    </template>
  </category>
  <category>
    <pattern>MOVIES</pattern>
```



```
<template>
YES <set name = "topic">MOVIES</set>
</template>
</category>
<category>
<pattern>HEALTH</pattern>
<template> YES <set name = "topic">HEALTH</set> </template>
</category>
<topic name ="MOVIES">
<category>
<pattern>*</pattern>
<template>
DO YOU LIKE COMEDY MOVIES?
</template>
</category>
<category> <pattern>YES</pattern>
<template>
I TOO LIKE COMEDY MOVIES
</template>
</category>
<category>
<pattern>NO</pattern>
<template>
BUT I LIKE COMEDY MOVIES
</template>
</category>
</topic>
<topic name ="HEALTH">
<category>
<pattern>*</pattern>
<template>
DO YOU HAVE FEVER?
</template>
</category>
<category>
<pattern>YES</pattern>
<template>
PLEASE TAKE MEDICINES AND PROPER REST
</template>
```

```
</category>
<category>
<pattern>NO</pattern>
<template>
GO OUT FOR A WALK AND LISTEN MUSIC
</template>
</category>
</topic>
  <category>
<pattern>NICE TALKING TO YOU</pattern>
<template>
SAME HERE...!!
</template>
</category>
</aiml>
```

Step 3: Install aiml packages

```
pip install aiml
pip install aimlbotkernel
or
pip3 install aiml
pip3 install aimlbotkernel
```

Step 4: Create chatbot.py file and run chatbot.py

```
import aiml
# Create the kernel and learn AIML files
kernel = aiml.Kernel()
kernel.learn("std-startup.xml")
kernel.respond("load aiml b")
# Press CTRL-C to break this loop
while True:
    message = input("Enter your message to the bot: ")
    if message == "quit":
        break
    else:
        bot_response = kernel.respond(message)
        print(bot_response)
```

OUTPUT:

```
Loading std-startup.xml...done (0.05 seconds)
Loading basic_chat.aiml...done (0.01 seconds)
Enter your message to the bot: Hello
WHAT WOULD YOU LIKE TO DISCUSS? : HEALTH, MOVIES
Enter your message to the bot: Health
YES HEALTH
Enter your message to the bot: I am feeling tired
DO YOU HAVE FEVER?
Enter your message to the bot: No
GO OUT FOR A WALK AND LISTEN MUSIC
Enter your message to the bot: Movies
YES MOVIES
Enter your message to the bot: I love movies
DO YOU LIKE COMEDY MOVIES?
Enter your message to the bot: Yes
I TOO LIKE COMEDY MOVIES
Enter your message to the bot: Nice talking to you
SAME HERE...!!
Enter your message to the bot: Quit
```

PRACTICAL3

AIM: Implement Bayes Theorem using Python.

CODE:

```
# calculate the probability of cancer patient and diagnostic test
# calculate P(A|B) given P(A), P(B|A), P(B|not A)
def bayes_theorem(p_a, p_b_given_a, p_b_given_not_a):
    # calculate P(not A)
    not_a = 1 - p_a
    # calculate P(B)
    p_b = p_b_given_a * p_a + p_b_given_not_a * not_a
    # calculate P(A|B)
    p_a_given_b = (p_b_given_a * p_a) / p_b
    return p_a_given_b

# P(A)
p_a = 0.0002
# P(B|A)
p_b_given_a = 0.85
# P(B|not A)
p_b_given_not_a = 0.05
# calculate P(A|B)
result = bayes_theorem(p_a, p_b_given_a, p_b_given_not_a)
# summarize
print('P(A|B) = %.3f%%' % (result * 100))
```

OUTPUT:

P(A|B) = 0.339%

PRACTICAL4

AIM: Implement Conditional Probability and joint probability using Python.

CODE:

```
import enum, random

class Kid(enum.Enum):

    BOY = 0

    GIRL = 1

def random_kid() -> Kid:

    return random.choice([Kid.BOY, Kid.GIRL])

both_girls = 0

older_girl = 0

either_girl = 0

random.seed(0)

for _ in range(10000):

    younger = random_kid()

    older = random_kid()

    if older == Kid.GIRL:

        older_girl += 1

    if older == Kid.GIRL and younger == Kid.GIRL:

        both_girls += 1

    if older == Kid.GIRL or younger == Kid.GIRL:

        either_girl += 1

print("older girl: ", older_girl)

print("both girl: ", both_girls)

print("either girl: ", either_girl)
```

```
print("P(both | older):", both_girls / older_girl)  
print("P(both | either):", both_girls / either_girl)
```

OUTPUT:

```
older girl:  4937  
both girl:   2472  
either girl: 7464  
P(both | older): 0.5007089325501317  
P(both | either): 0.3311897106109325
```

PRACTICAL5

AIM: Write a program for to implement Rule based system. (Prolog).

CODE:

```
go:-
    hypothesis(Disease),
    write('I believe that the patient have '),
    write(Disease),
    nl,
    write('TAKE CARE '),
    undo.
/*Hypothesis that should be tested*/
hypothesis(cold) :- cold, !.
hypothesis(flu) :- flu, !.
hypothesis(typhoid) :- typhoid, !.
hypothesis(measles) :- measles, !.
hypothesis(malaria) :- malaria, !.
hypothesis(unknown). /* no diagnosis*/
/*Hypothesis Identification Rules*/
cold :-
    verify(headache),
    verify(runny_nose),
    verify(sneezing),
    verify(sore_throat),
    write('Advices and Sugestions:'),
    nl,
    write('1: Tylenol/tab'),
    nl,
    write('2: panadol/tab'),
    nl,
    write('3: Nasal spray'),
    nl,
    write('Please wear warm cloths Because'),
    nl.
flu :-
    verify(fever),
```

```
verify(headache),
verify(chills),
verify(body_ache),
write('Advices and Sugestions:'),
nl,
write('1: Tamiflu/tab'),
nl,
write('2: panadol/tab'),
nl,
write('3: Zanamivir/tab'),
nl,
write('Please take a warm bath and do salt gargling Because'),
nl.
typhoid :-
verify(headache),
verify(abdominal_pain),
verify(poor_appetite),
verify(fever),
write('Advices and Sugestions:'),
nl,
write('1: Chloramphenicol/tab'),
nl,
write('2: Amoxicillin/tab'),
nl,
write('3: Ciprofloxacin/tab'),
nl,
write('4: Azithromycin/tab'),
nl,
write('Please do complete bed rest and take soft Diet Because'),
nl.
measles :-
verify(fever),
verify(runny_nose),
verify(rash),
verify(conjunctivitis),
write('Advices and Sugestions:'),
nl,
write('1: Tylenol/tab'),
nl,
```



```
write('2: Aleve/tab'),
nl,
write('3: Advil/tab'),
nl,
write('4: Vitamin A'),
nl,
write('Please Get rest and use more liquid Because'),
nl.
malaria :-
verify( fever ),
verify( sweating ),
verify( headache ),
verify( nausea ),
verify( vomiting ),
verify( diarrhea ),
write('Advices and Sugestions:'),
nl,
write('1: Aralen/tab'),
nl,
write('2: Qualaquin/tab'),
nl,
write('3: Plaquenil/tab'),
nl,
write('4: Mefloquine'),
nl,
write('Please do not sleep in open air and cover your full skin Because'),
nl.
/* how to ask questions */
ask( Question ) :-
write('Does the patient have following symptom:'),
write( Question ),
write('? '),
read( Response ),
nl,
( ( Response == yes ; Response == y )
->
assert( yes( Question ) ) ;
assert( no( Question ) ), fail ).
:- dynamic yes/1, no/1.
```

```

/*How to verify something */
verify(S) :-
  (yes(S)
  ->
  true ;
  (no(S)
  ->
  fail ;
  ask(S))).
/* undo all yes/no assertions*/
undo :- retract(yes(_)),fail.
undo :- retract(no(_)),fail.
undo.

```

OUTPUT:

```

?-
% c:/Users/PrateekKarkera/Downloads/daignosis (1).pl compiled 0.00 sec, 17 clauses
?- go.
Does the patient have following symptom:headache? yes.
Does the patient have following symptom:runny_nose? |: yes.
Does the patient have following symptom:sneezing? |: yes.
Does the patient have following symptom:sore_throat? |: yes.

Advices and Sugestions:
1: Tylenol/tab
2: panadol/tab
3: Nasal spray
Please wear warm cloths Because
I believe that the patient have cold
TAKE CARE
true.
?- ■

```

PRACTICAL 6

AIM: Design a Fuzzy based application using Python/ R.

CODE:

```
import numpy as np

import skfuzzy as fuzz

import matplotlib.pyplot as plt

from skfuzzy import control as ctrl

from mpl_toolkits.mplot3d import Axes3D # Required for 3D plotting


# New Antecedent/Consequent objects hold universe variables and membership
# functions


quality = ctrl.Antecedent(np.arange(0, 10, 0.1), 'quality')
service = ctrl.Antecedent(np.arange(0, 10, 0.1), 'service')
tip = ctrl.Consequent(np.arange(0, 25, 0.1), 'tip')


quality['poor'] = fuzz.zmf(quality.universe, 0,5)
quality['average'] = fuzz.gaussmf(quality.universe,5,1)
quality['good'] = fuzz.smf(quality.universe,5,10)


service['poor'] = fuzz.zmf(service.universe, 0,5)
service['average'] = fuzz.gaussmf(service.universe,5,1)
service['good'] = fuzz.smf(service.universe,5,10)


tip['low'] = fuzz.trimf(tip.universe, [0, 0, 13])
```

```
tip['medium'] = fuzz.trimf(tip.universe, [0, 13, 25])

tip['high'] = fuzz.trimf(tip.universe, [13, 25, 25])


quality['average'].view()

plt.title('Quality')


service['poor'].view()

plt.title('Service')


tip['medium'].view()

plt.title('Tip Medium')


rule1 = ctrl.Rule(quality['poor'] | service['poor'], tip['low'])
rule2 = ctrl.Rule(service['average'], tip['medium'])
rule3 = ctrl.Rule(service['good'] | quality['good'], tip['high'])

rule1.view()

plt.title('Rule 1')

rule2.view()

plt.title('Rule 2')

rule3.view()

plt.title('Rule 3')tipping_ctrl = ctrl.ControlSystem([rule1, rule2, rule3])
tipping = ctrl.ControlSystemSimulation(tipping_ctrl)


tipping.input['quality'] = 6.5
tipping.input['service'] = 9.8



tipping.compute()

print(tipping.output['tip'])


tip.view(sim=tipping)
```

```
plt.title('Result')  
plt.show(block=True)
```


OUTPUT:

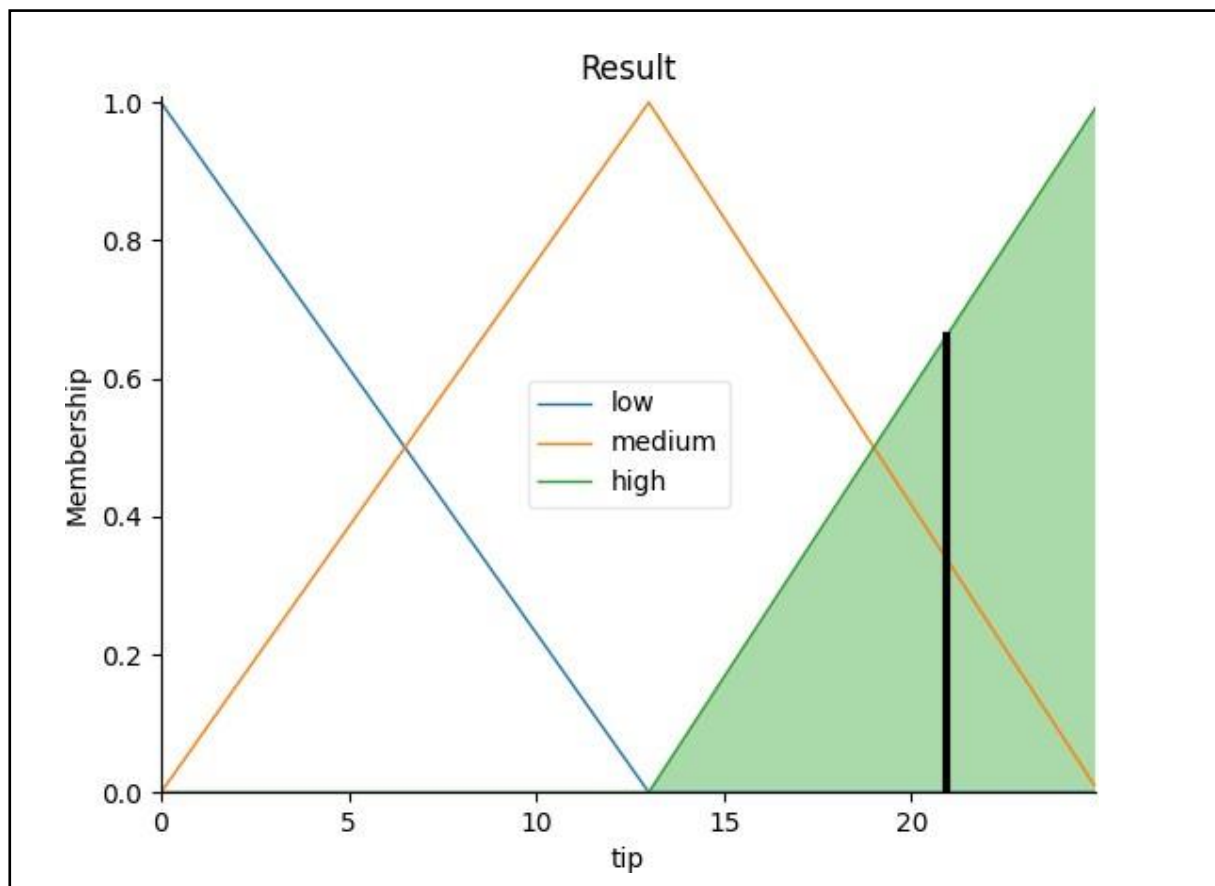



Rule 1



Rule 2





PRACTICAL7

[A] **AIM:** Write an application to stimulate supervised learning model.

CODE:

```
from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import classification_report, confusion_matrix

from sklearn import datasets

iris=datasets.load_iris()

x = iris.data

y = iris.target

print ('sepal-length', 'sepal-width', 'petal-length', 'petal-width')

print(x)

print('class: 0-Iris-Setosa, 1- Iris-Versicolour, 2- Iris-Virginica')

print(y)

x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)

#To Training the model and Nearest neighbors K=5

classifier = KNeighborsClassifier(n_neighbors=5)

classifier.fit(x_train, y_train)

#To make predictions on our test data

y_pred=classifier.predict(x_test)

print('Confusion Matrix')

print(confusion_matrix(y_test,y_pred))

print('Accuracy Metrics')

print(classification_report(y_test,y_pred))
```

OUTPUT:

[illegible]

[B] AIM: Write an application to stimulate unsupervised learning model.

CODE:

```
# Importing Modules
from scipy.cluster.hierarchy import linkage, dendrogram
import matplotlib.pyplot as plt
import pandas as pd

# Reading the DataFrame
seeds_df = pd.read_csv("seeds-less-rows.csv")

# Remove the grain species from the DataFrame, save for later
varieties = list(seeds_df.pop('grain_variety'))


# Extract the measurements as a NumPy array
samples = seeds_df.values

"""
Perform hierarchical clustering on samples using the
linkage() function with the method='complete' keyword argument.
Assign the result to mergings.
"""
mergings = linkage(samples, method='complete')

"""
Plot a dendrogram using the dendrogram() function on mergings,
specifying the keyword arguments labels=varieties, leaf_rotation=90,
and leaf_font_size=6.
"""
dendrogram(mergings,
            labels=varieties,
            leaf_rotation=90,
            leaf_font_size=6,
```

)

```
plt.show()
```

OUTPUT:

PRACTICAL 8

AIM: Write an application to implement clustering algorithm.

CODE:

```
import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.cluster import KMeans

import sklearn.metrics as sm

import pandas as pd

import numpy as np

iris = datasets.load_iris()

X = pd.DataFrame(iris.data)

X.columns = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width']

y = pd.DataFrame(iris.target)

y.columns = ['Targets']

model = KMeans(n_clusters=3)

model.fit(X)

plt.figure(figsize=(14,7))

colormap = np.array(['red', 'lime', 'black'])

# Plot the Original Classifications

plt.subplot(1, 2, 1)

plt.scatter(X.Petal_Length, X.Petal_Width,

c=colormap[y.Targets], s=40)

plt.title('Real Classification')

plt.xlabel('Petal Length')

plt.ylabel('Petal Width')
```

```

# Plot the Models Classifications

plt.subplot(1, 2, 2)

plt.scatter(X.Petal_Length, X.Petal_Width,
            c=colormap[model.labels_], s=40)

plt.title('K Mean Classification')

plt.xlabel('Petal Length')

plt.ylabel('Petal Width')

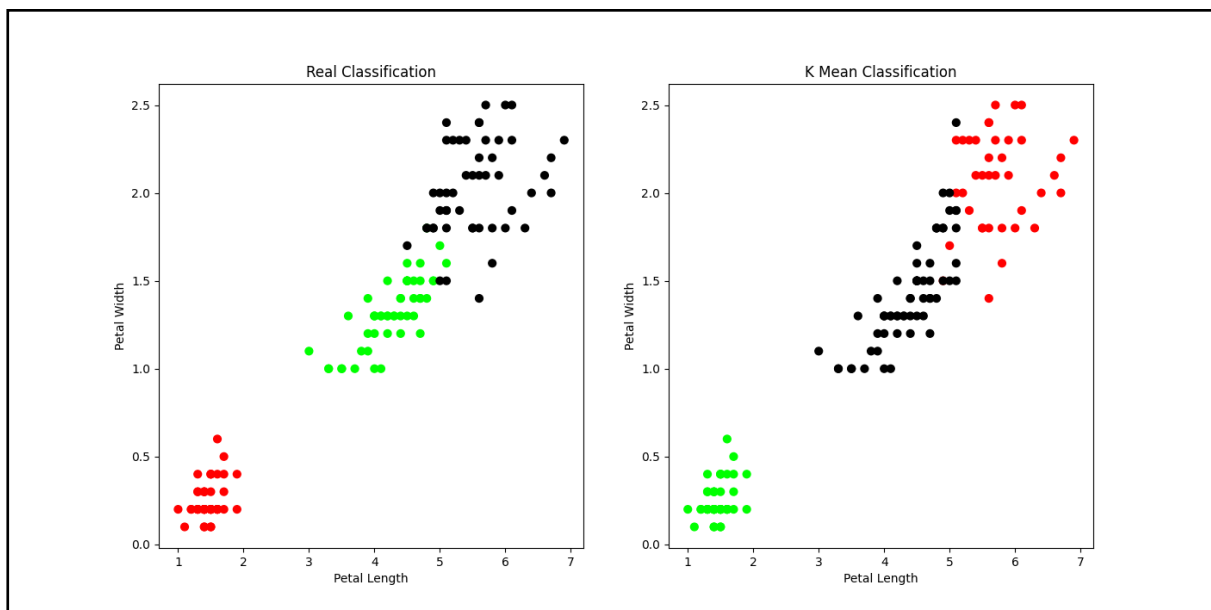
plt.show()

print('The accuracy score of K-Mean: ', sm.accuracy_score(y, model.labels_))

print('The Confusion matrix of K-Mean: ', sm.confusion_matrix(y,
model.labels ))

```

OUTPUT:



PRACTICAL9

AIM: Write an application to implement support vector machine algorithm.

CODE:

```
#Import scikit-learn dataset library

from sklearn import datasets


#Import svm model

from sklearn import svm


# Import train_test_split function

from sklearn.model_selection import train_test_split


#Import scikit-learn metrics module for accuracy calculation

from sklearn import metrics


#Load dataset

cancer = datasets.load_breast_cancer()


# print the names of the 13 features

print("Features: ", cancer.feature_names)


# print the label type of cancer('malignant' 'benign')

print("Labels: ", cancer.target_names)


# print data(feature) shape
```

```
cancer.data.shape

# print the cancer data features (top 5 records)
print(cancer.data[0:5])

# print the cancer labels (0:malignant, 1:benign)
print(cancer.target)

# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(cancer.data,
cancer.target, test_size=0.3,random_state=109) # 70% training and 30% test

#Create a svm Classifier
clf = svm.SVC(kernel='linear') # Linear Kernel

#Train the model using the training sets
clf.fit(X_train, y_train)

#Predict the response for test dataset
y_pred = clf.predict(X_test)

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))
```



```
# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:", metrics.recall_score(y_test, y_pred))
```

OUTPUT:

```
Features: ['mean radius' 'mean texture' 'mean perimeter' 'mean area'
'mean smoothness' 'mean compactness' 'mean concavity'
'mean concave points' 'mean symmetry' 'mean fractal dimension'
'radius error' 'texture error' 'perimeter error' 'area error'
'smoothness error' 'compactness error' 'concavity error'
'concave points error' 'symmetry error' 'fractal dimension error'
'worst radius' 'worst texture' 'worst perimeter' 'worst area'
'worst smoothness' 'worst compactness' 'worst concavity'
'worst concave points' 'worst symmetry' 'worst fractal dimension']
Labels: ['malignant' 'benign']
[[1.799e+01 1.038e+01 1.228e+02 1.001e+03 1.184e-01 2.776e-01 3.001e-01
 1.471e-01 2.419e-01 7.871e-02 1.095e+00 9.053e-01 8.589e+00 1.534e+02
 6.399e-03 4.904e-02 5.373e-02 1.587e-02 3.003e-02 6.193e-03 2.538e+01
 1.733e+01 1.846e+02 2.019e+03 1.622e-01 6.656e-01 7.119e-01 2.654e-01
 4.601e-01 1.189e-01]
[2.057e+01 1.777e+01 1.329e+02 1.326e+03 8.474e-02 7.864e-02 8.690e-02
 7.017e-02 1.812e-01 5.667e-02 5.435e-01 7.339e-01 3.398e+00 7.408e+01
 5.225e-03 1.308e-02 1.860e-02 1.340e-02 1.389e-02 3.532e-03 2.499e+01
 2.341e+01 1.588e+02 1.956e+03 1.238e-01 1.866e-01 2.416e-01 1.860e-01
 2.750e-01 8.902e-02]
[1.969e+01 2.125e+01 1.300e+02 1.203e+03 1.096e-01 1.599e-01 1.974e-01
 1.279e-01 2.069e-01 5.999e-02 7.456e-01 7.869e-01 4.585e+00 9.403e+01
 6.150e-03 4.006e-02 3.832e-02 2.058e-02 2.250e-02 4.571e-03 2.357e+01
 2.553e+01 1.525e+02 1.709e+03 1.444e-01 4.245e-01 4.504e-01 2.430e-01
 3.613e-01 8.758e-02]
[1.142e+01 2.038e+01 7.758e+01 3.861e+02 1.425e-01 2.839e-01 2.414e-01
 1.052e-01 2.597e-01 9.744e-02 4.956e-01 1.156e+00 3.445e+00 2.723e+01
 9.110e-03 7.458e-02 5.661e-02 1.867e-02 5.963e-02 9.208e-03 1.491e+01
 2.650e+01 9.887e+01 5.677e+02 2.098e-01 8.663e-01 6.869e-01 2.575e-01
 6.638e-01 1.730e-01]
[2.029e+01 1.434e+01 1.351e+02 1.297e+03 1.003e-01 1.328e-01 1.980e-01
 1.043e-01 1.809e-01 5.883e-02 7.572e-01 7.813e-01 5.438e+00 9.444e+01
 1.149e-02 2.461e-02 5.688e-02 1.885e-02 1.756e-02 5.115e-03 2.254e+01
 1.667e+01 1.522e+02 1.575e+03 1.374e-01 2.050e-01 4.000e-01 1.625e-01
 2.364e-01 7.678e-02]]
```

[illegible]

Accuracy: 0.9649122807017544

Precision: 0.9811320754716981

Recall: 0.9629629629629629

PRACTICAL 10

AIM: Simulate artificial neural network model with both feedforward and backpropagation approach.

CODE:

```
import numpy as np

X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float) # two inputs
[sleep, study]

y = np.array([92, 86, 89], dtype=float) # one output [Expected % in Exams]

X = X / np.amax(X, axis=0) # maximum of X array longitudinally

y = y / 100

# Sigmoid Function

def sigmoid(x):

    return 1 / (1 + np.exp(-x))

# Derivative of Sigmoid Function

def derivatives_sigmoid(x):

    return x * (1 - x)

# Variable initialization

epoch = 5000 # Setting training iterations

lr = 0.1 # Setting learning rate

inputlayer_neurons = 2 # number of features in data set

hiddenlayer_neurons = 3 # number of hidden layers neurons

output_neurons = 1 # number of neurons at output layer
```

```

# weight and bias initialization

wh = np.random.uniform(size=(inputlayer_neurons, hiddenlayer_neurons)) #
weight of the link from input node to hidden node

bh = np.random.uniform(size=(1, hiddenlayer_neurons)) # bias of the link
from input node to hidden node

wout = np.random.uniform(size=(hiddenlayer_neurons, output_neurons)) #
weight of the link from hidden node to output node

bout = np.random.uniform(size=(1, output_neurons)) # bias of the link from
hidden node to output node

# draws a random range of numbers uniformly of dim x*y

for i in range(epoch):

    # Forward Propagation

    hinp1 = np.dot(X, wh)

    hinp = hinp1 + bh

    hlayer_act = sigmoid(hinp)

    outinp1 = np.dot(hlayer_act, wout)

    outinp = outinp1 + bout

    output = sigmoid(outinp)

    # Backpropagation

    EO = y - output

    outgrad = derivatives_sigmoid(output)

    d_output = EO * outgrad

    EH = d_output.dot(wout.T)

    # how much hidden layer weights contributed to error

    hiddengrad = derivatives_sigmoid(hlayer_act)

    d_hiddenlayer = EH * hiddengrad

# dotproduct of nextlayererror and currentlayerop

wout += hlayer_act.T.dot(d_output) * lr

```

```
wh += X.T.dot(d_hiddenlayer) * lr  
  
print("Input: \n" + str(X))  
  
print("Actual Output: \n" + str(y))  
  
print("Predicted Output: \n", output)
```

OUTPUT:

```
Input:  
[[0.66666667 1.          ]  
 [0.33333333 0.55555556]  
 [1.          0.66666667]]  
Actual Output:  
[[0.92]  
 [0.86]  
 [0.89]]  
Predicted Output:  
[[0.84047843]  
 [0.81670721]  
 [0.83471132]]
```



Parle Tilak Vidyalaya Association's
MULUND COLLEGE OF COMMERCE
(AUTONOMOUS)

(Affiliated to University of Mumbai)
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that SONAWANE CHAITANYA RAJESH of **M.Sc.**
I.T. Part II Roll No 2414559 has successfully completed the practical
work in Machine Learning in partial fulfilment of the requirements for
the Semester III of **M.Sc. I.T. Part II** during the academic year
2024-25.

Teacher In-charge and Coordinator

Examiner

Date:

College Seal

INDEX

Practical No.	Name of Practical	Page No.	Signature
1	1.1 Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers 1.2 Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables 1.3 Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.	3	
2	2.1 Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a CSV file and generate the final specific hypothesis. (Create your dataset)	15	
3	3.1 Simple Linear Regression Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE 3.2 Multiple Linear Regression Extend linear regression to multiple features. Handle feature selection and potential multicollinearity 3.3 Regularized Linear Models (Ridge, Lasso, ElasticNet) Implement regression variants like LASSO and Ridge on any generated dataset.	16	
4	4.1 Logistic Regression Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the	28	

	ROC curve. 4.2 Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions		
5	5.1 Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample. 5.2 Implement Hidden Markov Models using hmmlearn	44	
6	6.1 Implement Bayesian Linear Regression to explore prior and posterior distribution 6.2 Implement Gaussian Mixture Models for density estimation and unsupervised clustering	49	
7	7.1 Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation 7.2 Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)	51	
8	Implement Bayesian Learning using inferences	56	
9	Set up a generator network to produce samples and a discriminator network to distinguish between real and generated data. (Use a simple small dataset	57	
10	Develop an API to deploy your model and perform predictions	61	

Practical 1

1.1 Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.

```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
plt.rcParams["figure.figsize"] = [8,6]
sns.set(style="darkgrid")
```

```
df = sns.load_dataset('titanic')
```

```
print(df.head())
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class
0 0	3	male	22.0	1	0	7.2500	S	Third	
1 1	1	female	38.0	1	0	71.2833	C	First	
2 1	3	female	26.0	0	0	7.9250	S	Third	
3 1	1	female	35.0	1	0	53.1000	S	First	
4 0	3	male	35.0	0	0	8.0500	S	Third	

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
print(df.isnull().sum())
```

```
survived    0
pclass      0
sex         0
age        177
sibsp       0
parch       0
fare        0
embarked    2
class       0
who         0
adult_male  0
deck       688
embark_town 2
alive       0
```

```
alone      0
dtype: int64
```

```
df      =      df[["age",      "embarked"]]
print(df.head())
```

```
      age embarked
0  22.0      S
1  38.0      C
2  26.0      S
3  35.0      S
4  35.0      S
```

```
df.loc[:, 'age'] = df.age.fillna(df.age.median())
df = df.dropna(subset=["embarked"])
```

```
print(df.head(20))
```

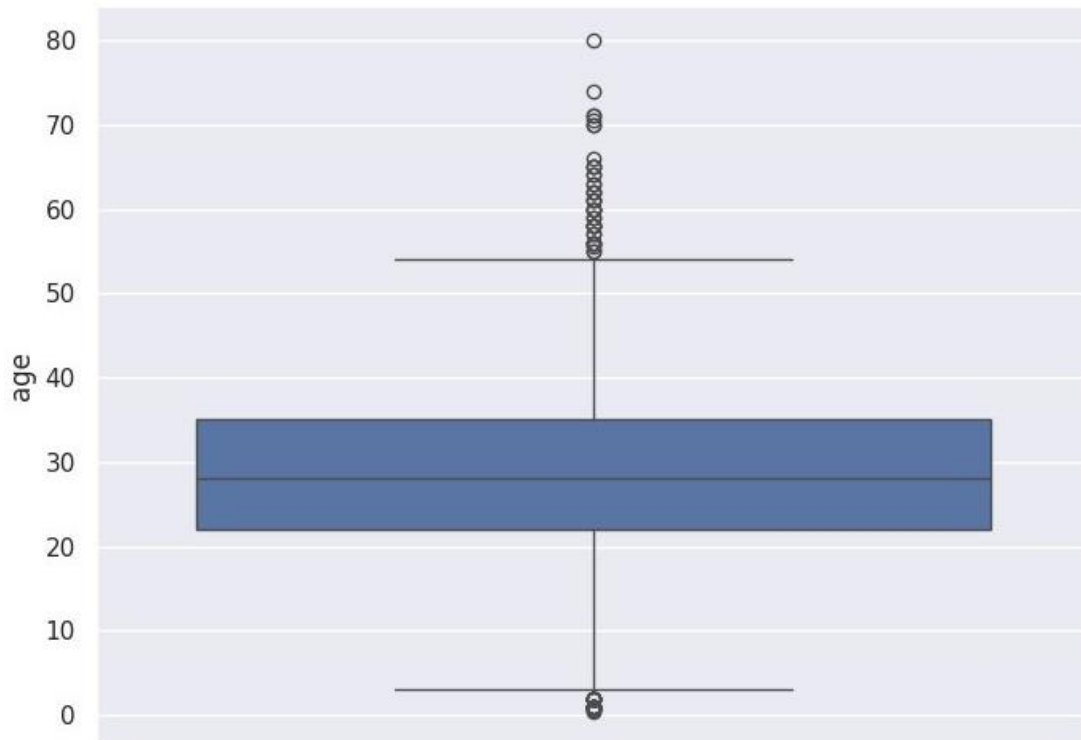
```
      age embarked
0  22.0      S
1  38.0      C
2  26.0      S
3  35.0      S
4  35.0      S
5  28.0      Q
6  54.0      S
7   2.0      S
8  27.0      S
9  14.0      C
10  4.0      S
11 58.0      S
12 20.0      S
13 39.0      S
14 14.0      S
15 55.0      S
16  2.0      Q
17 28.0      S
18 31.0      S
19 28.0      C
```

```
df.loc[:, 'embarked'] = df.embarked.str.upper()
print(df.embarked.unique())
```

```
['S' 'C' 'Q']
```

```
sns.boxplot(data=df.age)
```

```
<Axes: ylabel='age'>
```



```
Q1 = df.age.quantile(0.25)
Q3 = df.age.quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

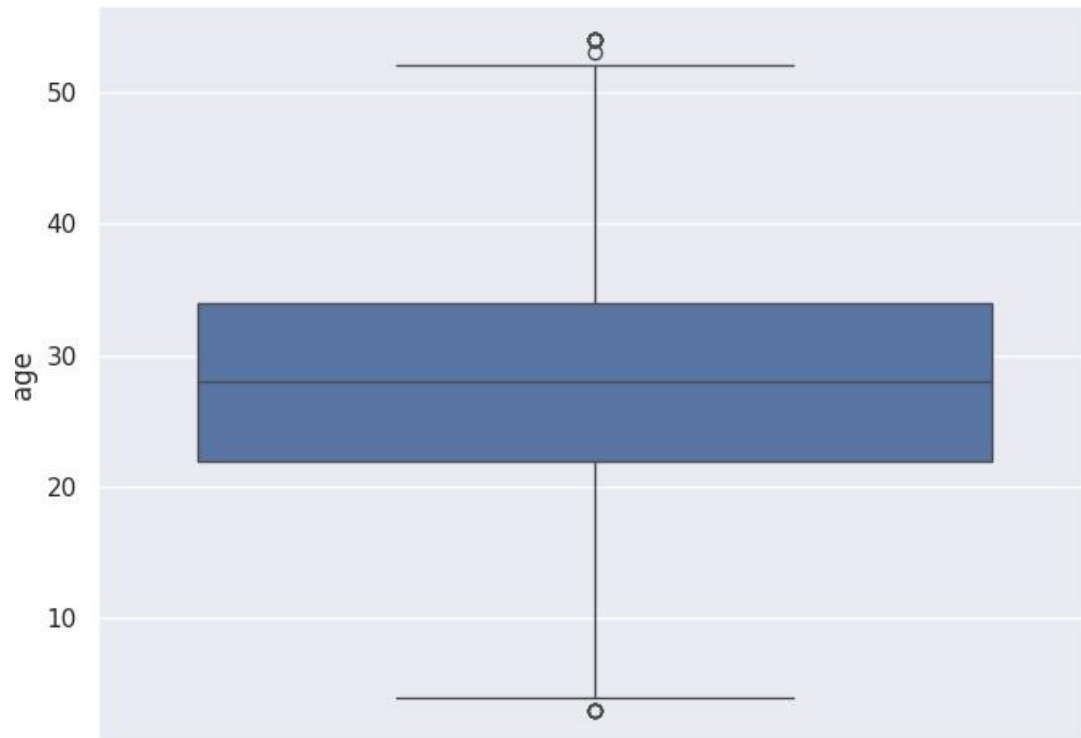
df = df[(df.age >= lower_bound) & (df.age <= upper_bound)]

print(df.head())
```

	age	embarked
0	22.0	S
1	38.0	C
2	26.0	S
3	35.0	S
4	35.0	S

```
sns.boxplot(data=df.age)

<Axes: ylabel='age'>
```



1.2 Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

plt.rcParams["figure.figsize"] = [8,6]
sns.set(style="darkgrid")

df = sns.load_dataset('iris')
print(df.head())
```

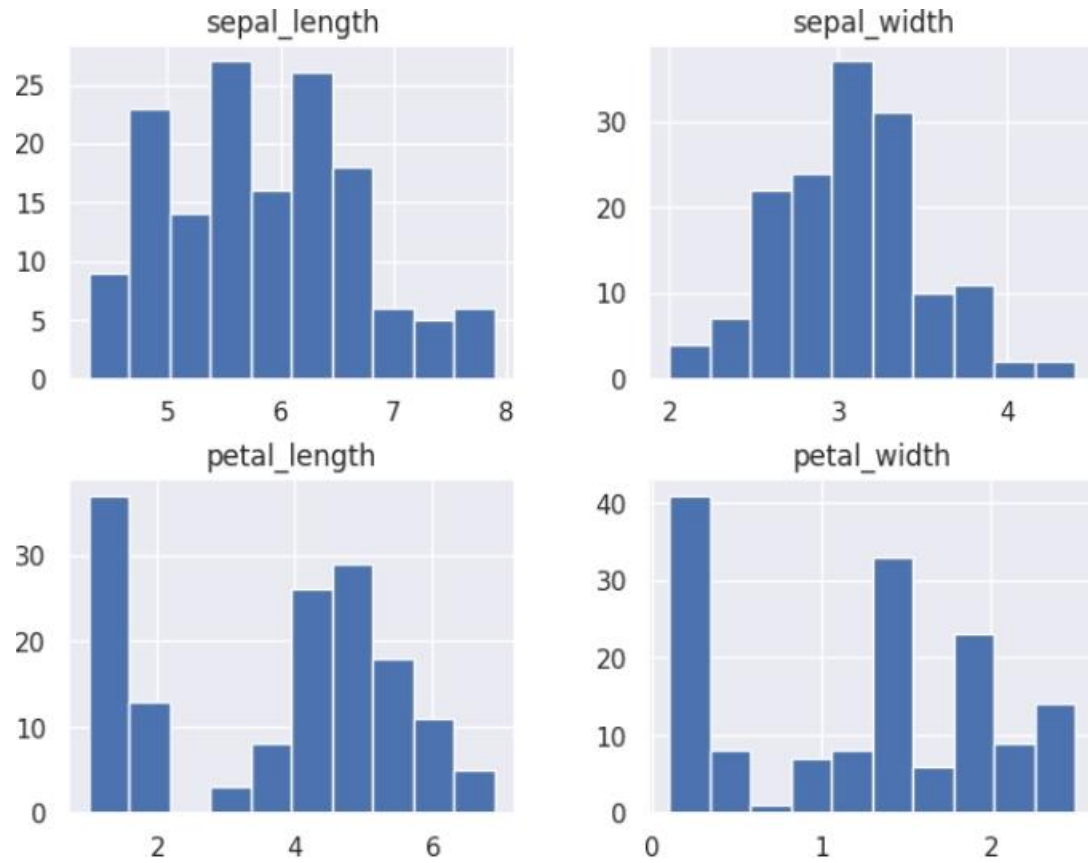
	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
summary_statistics = df.describe()
print(summary_statistics)
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000


```
#Univariate Visualizations
df.hist()

array([[<Axes: title={'center': 'sepal_length'}>,
        <Axes: title={'center': 'sepal_width'}>],
       [<Axes: title={'center': 'petal_length'}>,
        <Axes: title={'center': 'petal_width'}>]], dtype=object)
```




```
sns.boxplot(data=df)
```

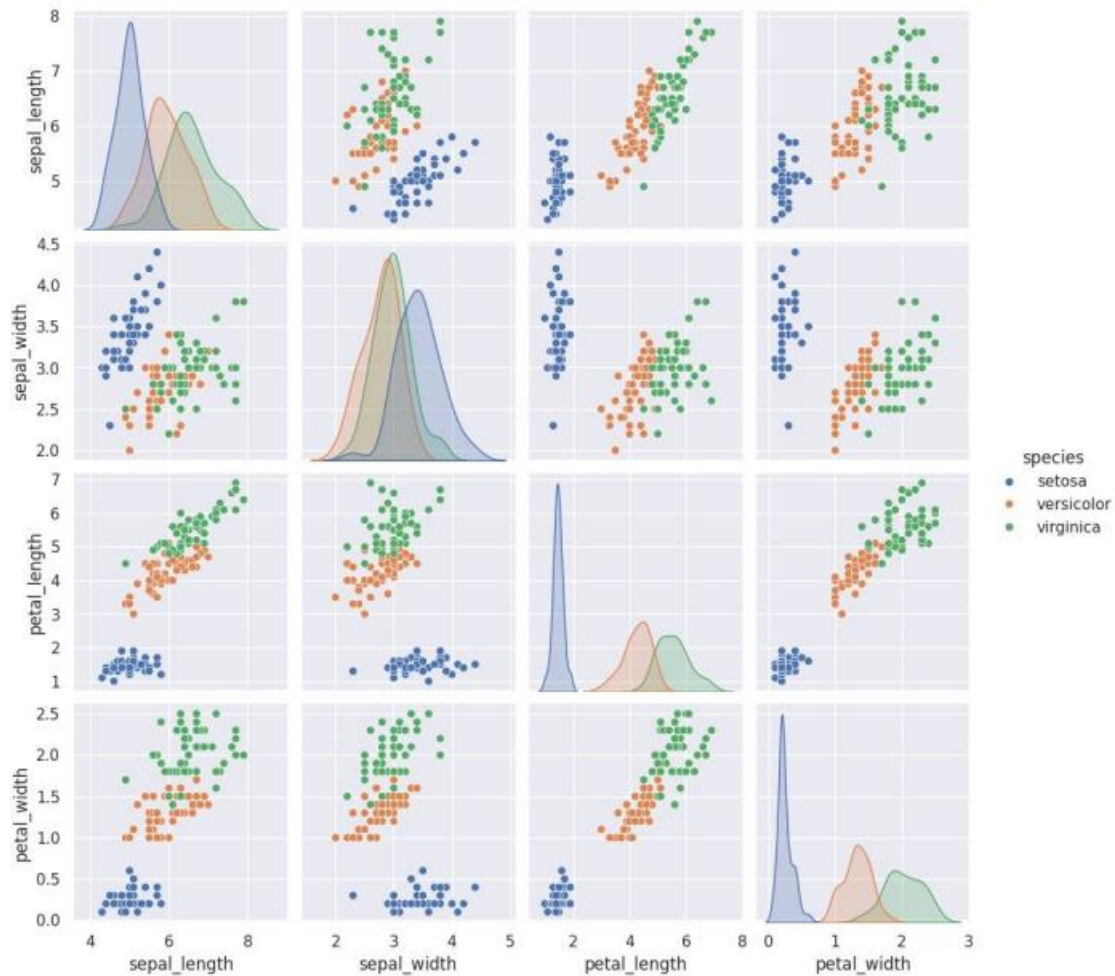
<Axes: >



```
sns.scatterplot(data=df, x='sepal_length', y='sepal_width', hue='species')  
<Axes: xlabel='sepal_length', ylabel='sepal_width'>
```



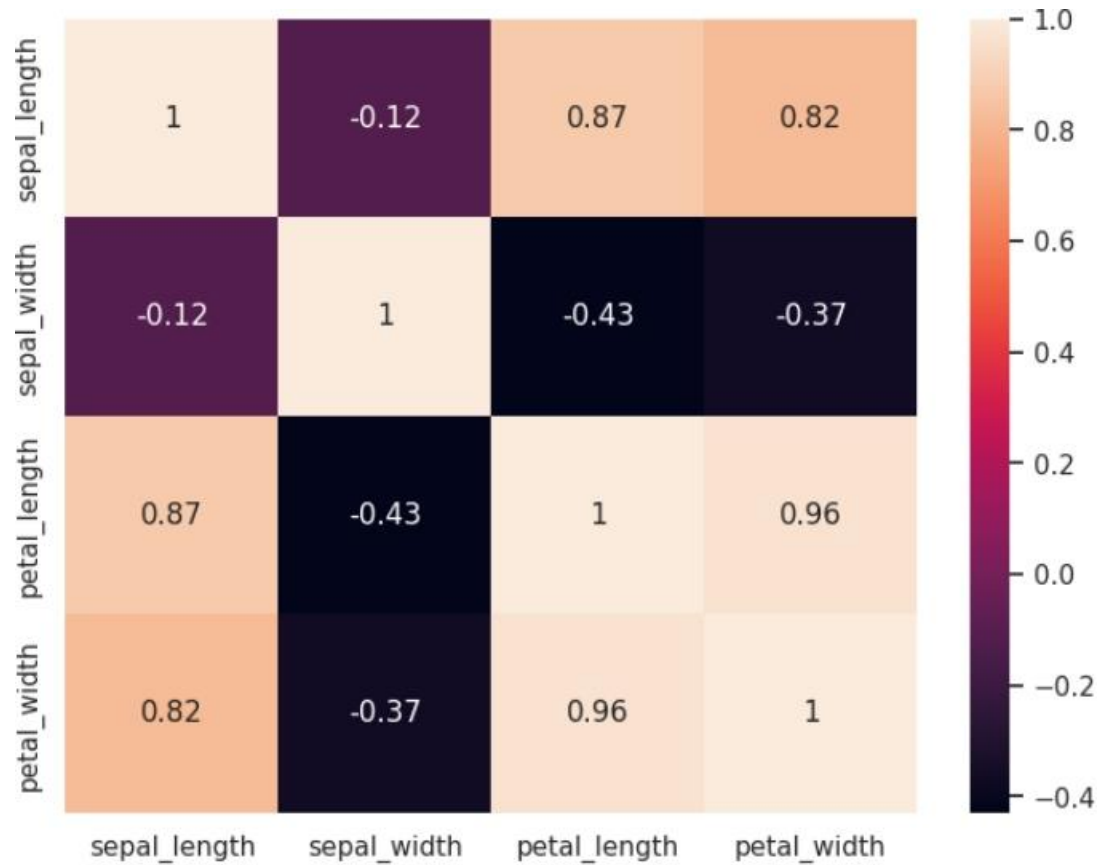
```
sns.pairplot(df, hue="species")  
<seaborn.axisgrid.PairGrid at 0x7fc52c528610>
```

#Correlation

```
numeric_df = df.select_dtypes(include=[np.number])
correlation_matrix = numeric_df.corr()
sns.heatmap(correlation_matrix, annot=True)
```

<Axes: >



```
potential_features = df.select_dtypes(include=[np.number]).columns.tolist()
target_variable = 'species'
```

```
print("Potential Features: ", potential_features)
print("Target Variable: ", target_variable)
```

```
Potential Features: ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
Target Variable: species
```

1.3 Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.

```
import seaborn as sns
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import Binarizer
```

```
df = sns.load_dataset('titanic')
print(df.head())
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class
0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
label_encoder = LabelEncoder()
0 1
1 0
2 0
3 0
4 1
df['sex'] = label_encoder.fit_transform(df['sex'])
print(df.sex.head())
```

Name: sex, dtype: int64

```
scaler = StandardScaler()
df[['age', 'fare']] = scaler.fit_transform(df[['age', 'fare']])
print(df[['age', 'fare']].head())
```

	age	fare
0	-0.530377	-0.502445
1	0.571831	0.786845
2	-0.254825	-0.488854

```
3 0.365167 0.420730
4 0.365167 -0.486337
```

```
binarizer = Binarizer(threshold=0)
df[['fare']] = binarizer.fit_transform(df[['fare']])
print(df.fare.head())
```

```
0 0.0
1 1.0
2 0.0
3 1.0
4 0.0
```

```
Name: fare, dtype: float6
```

Practical 2

2.1 Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a CSV file and generate the final specific hypothesis. (Create your dataset)

```
import pandas as pd

data = {
    'Material': ['Plastic', 'Metal', 'Glass', 'Metal'],
    'Color': ['White', 'Silver', 'Green', 'Grey'],
    'Size': ['Small', 'Large', 'Small', 'Large'],
    'Recyclable': ['Yes', 'Yes', 'Yes', 'No'],
    'E-Waste': ['No', 'Yes', 'No', 'Yes']
}

df = pd.DataFrame(data)
df.to_csv('training_data.csv', index=False)

data = pd.read_csv('training_data.csv')
X = data.iloc[:, :-1]
y = data.iloc[:, -1]

hypothesis = ['0'] * X.shape[1]

for i in range(len(X)):
    if y[i] == 'Yes':
        for j in range(X.shape[1]):
            if hypothesis[j] == '0':
                hypothesis[j] = X.iloc[i, j]
            elif hypothesis[j] != X.iloc[i, j]:
                hypothesis[j] = '?'

print(hypothesis)

['Metal', '?', 'Large', '?']
```

Practical 3

3.1 Simple Linear Regression Fit a linear regression model on a dataset.

Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE


```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

np.random.seed(42)
X = 2.5 * np.random.rand(100, 1)
y = 5 + 2 * X + np.random.randn(100, 1)

data = pd.DataFrame({'Feature': X.flatten(), 'Target': y.flatten()})
print(data.head())
```

	Feature	Target
0	0.936350	6.959748
1	2.376786	9.454564
2	1.829985	8.751730
3	1.496646	6.005724
4	0.390047	5.560421

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Feature', y='Target', data=data)
plt.title('Feature vs Target')
plt.show()
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
andom_state=42)
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
print(f"Intercept: {model.intercept_[0]:.2f}")
```

```
print(f"Coefficient: {model.coef_[0][0]:.2f}")
```

```
Intercept: 5.14
```

```
Coefficient: 1.84
```

```
y_pred = model.predict(X_test)
```

```
pred_df = pd.DataFrame({'Actual': y_test.flatten(), 'Predicted': y_pred.fl
atten()})
```

```
print(pred_df.head())
```

	Actual	Predicted
0	5.974345	5.435196
1	8.970661	9.257909
2	7.624273	8.694195


```
3 7.4032248.189620
4 7.0849326.332951
```

```
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.2f}")
```

```
r2 = r2_score(y_test, y_pred)
print(f"R-squared: {r2:.2f}")
```

```
Mean Squared Error (MSE): 0.65
R-squared: 0.73
```

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Feature', y='Target', data=data, label='Actual Data')
plt.plot(X_test, y_pred, color='red', label='Regression Line')
plt.title('Linear Regression Fit')
plt.legend()
plt.show()
```



3.2 Multiple Linear Regression Extend linear regression to multiple features.

Handle feature selection and potential multicollinearity.

```
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor


df = sns.load_dataset('diamonds')

print("Missing values in the dataset:")
print(df.isnull().sum())

df = pd.get_dummies(df, drop_first=True)

Missing values in the dataset:
carat      0
cut         0
color      0
clarity     0
depth       0
table       0
price       0
x           0
y           0
z           0
dtype: int64

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()
```



```
X = df[['carat', 'depth', 'table', 'x', 'y', 'z',
        'cut_Premium', 'cut_Good', 'cut_Very Good',
        'color_E', 'color_F', 'clarity_VS2', 'clarity_VS1']]
y = df['price']

y = y.astype(float)

X_with_constant = sm.add_constant(X)

X_with_constant = X_with_constant.astype(int)

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X_with_constant.values, i+1) for i
in range(len(X.columns))]
print(vif)
```

	Features	VIF
0	carat	3.613538

```

2 table 1.530672
3 x 19.224267
4 y 15.677513
5 z 5.789510
6 cut_Premium 1.548643
7 cut_Good 1.295429
8 cut_Very Good 1.346362
9 color_E 1.079848
10 color_F 1.060367
11 clarity_VVS2 1.051655
12 clarity_VS1 1.031049

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, r
andom_state=42)

```

```

print("Data types of X_train:")
print(X_train.dtypes)
print("Data type of y_train:", y_train.dtype)

```

Data types of X_train:

```

carat    float64
depth    float64
table    float64
x         float64
y         float64
z         float64

```

```

cut_Premium    bool
cut_Good        bool
cut_Very Good  bool
color_E        bool
color_F        bool
clarity_VVS2   bool
clarity_VS1    bool
dtype: object

```

Data type of y_train: float64

```

X_train = X_train.astype(float)
y_train = y_train.astype(float)

```

```

model = LinearRegression()
model.fit(X_train, y_train)

```

```

print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)

```

Intercept: 17520.480548853404

Coefficients: [1.06799558e+04 -1.74848962e+02 -8.87411825e+01 -1.17618393e+03

3.03543071e+01 8.16330490e+00 -3.94552416e+01 -1.98436785e+02

```
-1.87877044e+01    4.36586414e+02    4.74099129e+02    1.02842658e+03
6.62232418e+02]
```

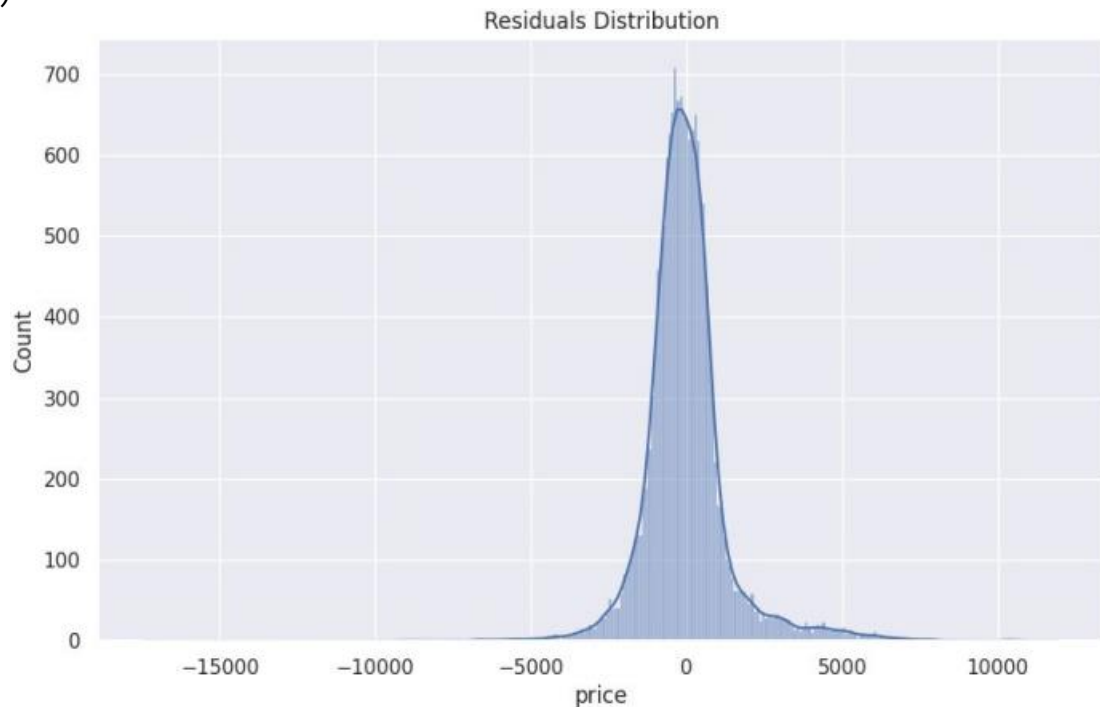
```
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
print("R-squared:", r2)
```


```
residuals = y_test - y_pred
```

```
Mean Squared Error: 2018911.748442661
R-squared: 0.8705490836162249
```

```
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.title("Residuals Distribution")
plt.show()
```



```
plt.scatter(y_pred, residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residuals vs Predicted Values")
plt.show()
```



```
X_train_sm = sm.add_constant(X_train)
ols_model = sm.OLS(y_train, X_train_sm).fit()
print(ols_model.summary())
```

OLS Regression Results

```
=====
===
Dep. Variable:      price R-squared:  0
.870
Model: OLS      Adj. R-squared:  0
.869
Method: Least Squares      F-statistic:1.935
e+04
Date:   Thu, 24 Oct 2024   Prob (F-statistic): 0.00
Time:   07:44:35         Log-Likelihood:   -3.2835 e+05
No. Observations:  37758 AIC:   6.567
e+05
Df Residuals: 37744 BIC:   6.569
e+05
Df Model:      13
Covariance Type:  nonrobust
=====
=====
```

```

                                coef      std err      t      P>|t| [0.025
0.975]
-----
const      1.752e+04    537.976    32.567 0.000 1.65e+04    1
.86e+04
carat      1.068e+04    72.935 146.431    0.000 1.05e+04    1
.08e+04
depth      -174.8490    6.336 -27.595    0.000 -187.268    -
162.430
table      -88.7412    4.144 -21.414    0.000 -96.864
-80.619
x      -1176.1839    46.936 -25.059    0.000 -1268.179    -1
084.188
y      30.3543    27.897 1.088 0.277 -24.325
85.034
z      8.16334 3.889 0.186 0.852 -77.861
94.188
cut_Premium      -39.4552    21.233 -1.858 0.063 -81.073
2.163
cut_Good      -198.4368    29.625 -6.698 0.000 -256.502    -
140.372
cut_Very Good      -18.7877    20.712 -0.907 0.364 -59.383
21.808
color_E      436.5864    20.108 21.712 0.000 397.174
475.999
color_F      474.0991    20.084 23.606 0.000 434.735
513.463
clarity_VVS2      1028.4266    26.287 39.123 0.000 976.903    1
079.950
clarity_VS1      662.2324    21.075 31.423 0.000 620.925
703.540
=====
====
Omnibus:    9105.593    Durbin-Watson:    1
.992
Prob(Omnibus):    0.000 Jarque-Bera (JB): 327369
.326
Skew:      0.453 Prob(JB):
0.00
Kurtosis: 17.397 Cond. No.    6.14
e+03
=====
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
The condition number is large, 6.14e+03. This might indicate that there are
strong multicollinearity or other numerical problems.

```

3.3 Multiple Linear Regression Extend linear regression to multiple features.

Handle feature selection and potential multicollinearity.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.metrics import mean_squared_error

X, y, coef = make_regression(n_samples=100, n_features=10, noise=0.1, coef
                             =True, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
                                                    andom_state=42)

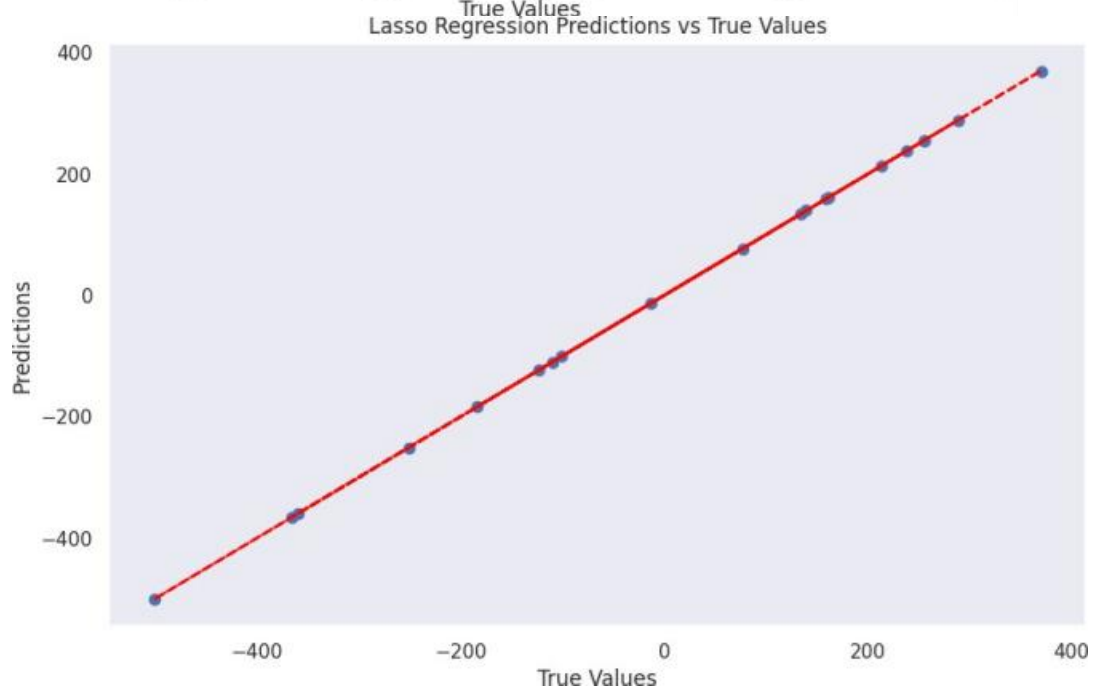
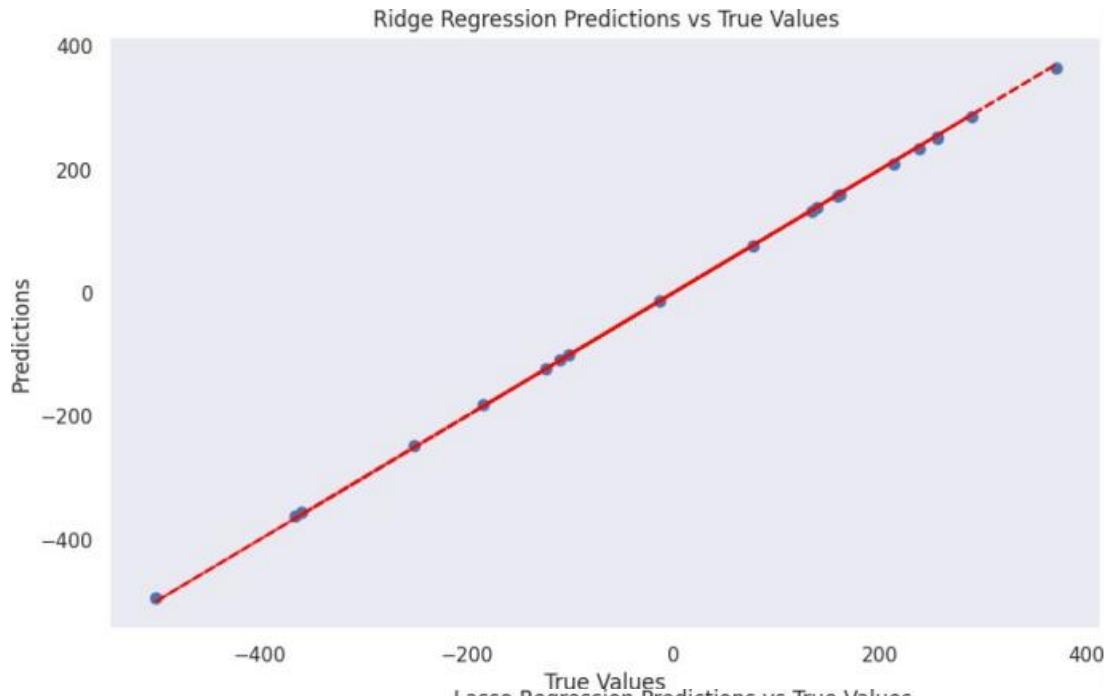
ridge_model = Ridge(alpha=1.0)
ridge_model.fit(X_train, y_train)
ridge_pred = ridge_model.predict(X_test)


lasso_model = Lasso(alpha=0.1)
lasso_model.fit(X_train, y_train)
lasso_pred = lasso_model.predict(X_test)

elastic_model = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_model.fit(X_train, y_train)
elastic_pred = elastic_model.predict(X_test)

def plot_results(y_test, predictions, model_name):
    plt.figure(figsize=(10, 6))
    plt.scatter(y_test, predictions)
    plt.plot(y_test, y_test, color='red', linestyle='--') # y=x line
    plt.title(f'{model_name} Predictions vs True Values')
    plt.xlabel('True Values')
    plt.ylabel('Predictions')
    plt.grid()
    plt.show()

plot_results(y_test, ridge_pred, 'Ridge Regression')
plot_results(y_test, lasso_pred, 'Lasso Regression')
plot_results(y_test, elastic_pred, 'ElasticNet Regression')
```





```
print("Mean Squared Error (MSE):")
print(f"Ridge:      {mean_squared_error(y_test,      ridge_pred):.2f}")
print(f"Lasso:      {mean_squared_error(y_test,      lasso_pred):.2f}")
print(f"ElasticNet: {mean_squared_error(y_test, elastic_pred):.2f}")
```

Mean Squared Error (MSE):
Ridge: 11.84
Lasso: 0.18
ElasticNet: 176.03

Practical 4

4.1 Logistic Regression Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score,
roc_curve, auc
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
logistic_reg_model = LogisticRegression()
logistic_reg_model.fit(X_train, y_train)
y_pred = logistic_reg_model.predict(X_test)


accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")

Accuracy: 0.85
Precision: 0.89
Recall: 0.82

y_prob = logistic_reg_model.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='blue', lw=2, label=f"ROC Curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], color='gray', linestyle='--') # Diagonal line for random classifier
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
```

```
plt.legend(loc='lower right')  
plt.show()
```



4.2 Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.

```

from sklearn import datasets
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report

iris = datasets.load_iris()
X = iris.data
y = iris.target

df = pd.DataFrame(data=X, columns=iris.feature_names)
df['target'] = y

```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.
2				
1	4.9	3.0	1.4	0.
2				
2	4.7	3.2	1.3	0.
2				
3	4.6	3.1	1.5	0.
2				
4	5.0	3.6	1.4	0.
2				

```
print(df.head())
```

```

sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
) \
0      5.1          3.5          1.4          0.
2
1      4.9          3.0          1.4          0.
2
2      4.7          3.2          1.3          0.
2
3      4.6          3.1          1.5          0.
2
4      5.0          3.6          1.4          0.

```

```
print(df.isnull().sum())
```

```
sepal length (cm)    0
sepal width (cm)     0
petal length (cm)    0
petal width (cm)     0
target              0
dtype: int64
```

```
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
y = y.astype('category').cat.codes
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
andom_state=42)
```

```
k = 3
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
```

```
y_pred = knn.predict(X_test)
```

```
for i in range(len(y_test)):
    print(f'Predicted: {iris.target_names[y_pred[i]]}, Actual: {iris.targe
t_names[y_test.iloc[i]]}')
```

```
Predicted: versicolor, Actual: versicolor Predicted: setosa, Actual: setosa
Predicted: virginica, Actual: virginica Predicted: versicolor, Actual:
versicolor Predicted: versicolor, Actual: versicolor Predicted: setosa,
Actual: setosa Predicted: versicolor, Actual: versicolor Predicted:
virginica, Actual: virginica Predicted: versicolor, Actual: versicolor
Predicted: versicolor, Actual: versicolor Predicted: virginica, Actual:
virginica Predicted: setosa, Actual: setosa Predicted: setosa, Actual:
setosa Predicted: setosa, Actual: setosa Predicted: setosa, Actual: setosa
```

```

Predicted:  setosa,    Actual:  setosa
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted: versicolor, Actual: versicolor
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted:  setosa,    Actual:  setosa
Predicted: virginica, Actual: virginica
Predicted:  setosa,    Actual:  setosa
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted:  setosa,    Actual:  setosa
Predicted: setosa, Actual: setosa

```

```

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')

```

```

print(classification_report(y_test, y_pred, target_names=iris.target_names
))

```

Accuracy: 100.00%

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

4.3 Build a decision tree classifier or regressor. Control hyperparameters

like tree depth to avoid overfitting. Visualize the tree.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris, fetch_california_housing
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, mean_squared_error

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

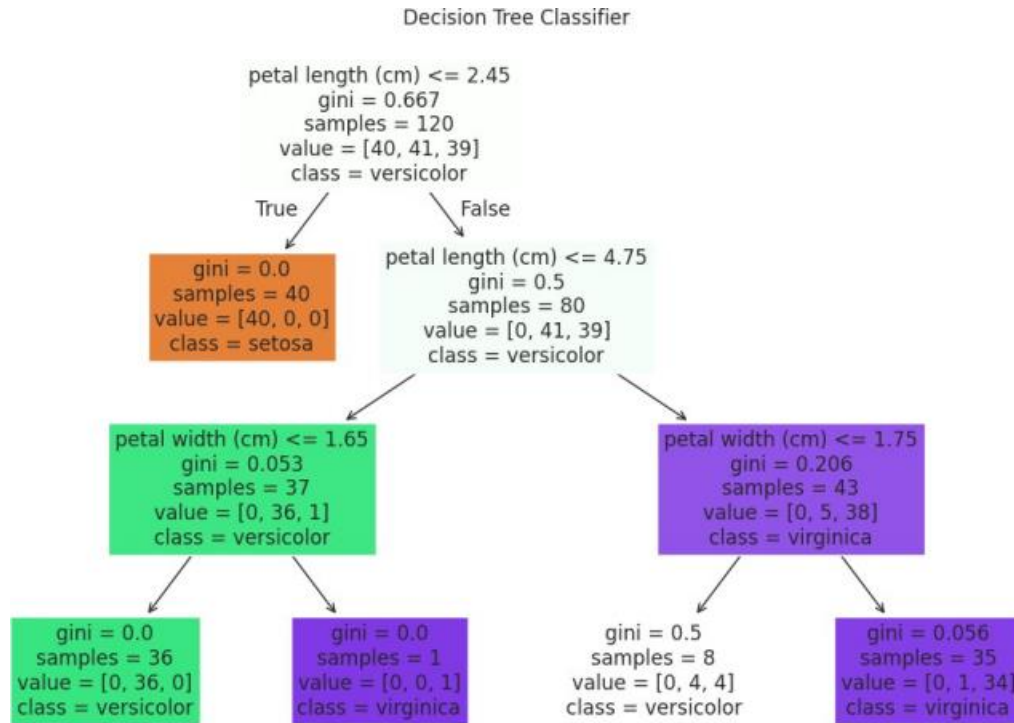
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')

Accuracy:          1.00

plt.figure(figsize=(12,8))
plot_tree(clf, filled=True, feature_names=iris.feature_names, class_names=iris.target_names)
plt.title("Decision Tree Classifier")
plt.show()
```



```

housing = fetch_california_housing()
X_housing = housing.data
y_housing = housing.target

X_train_housing, X_test_housing, y_train_housing, y_test_housing = train_test_split(X_housing, y_housing, test_size=0.2, random_state=42)

reg = DecisionTreeRegressor(max_depth=3, random_state=42)
reg.fit(X_train_housing, y_train_housing)

y_pred_housing = reg.predict(X_test_housing)

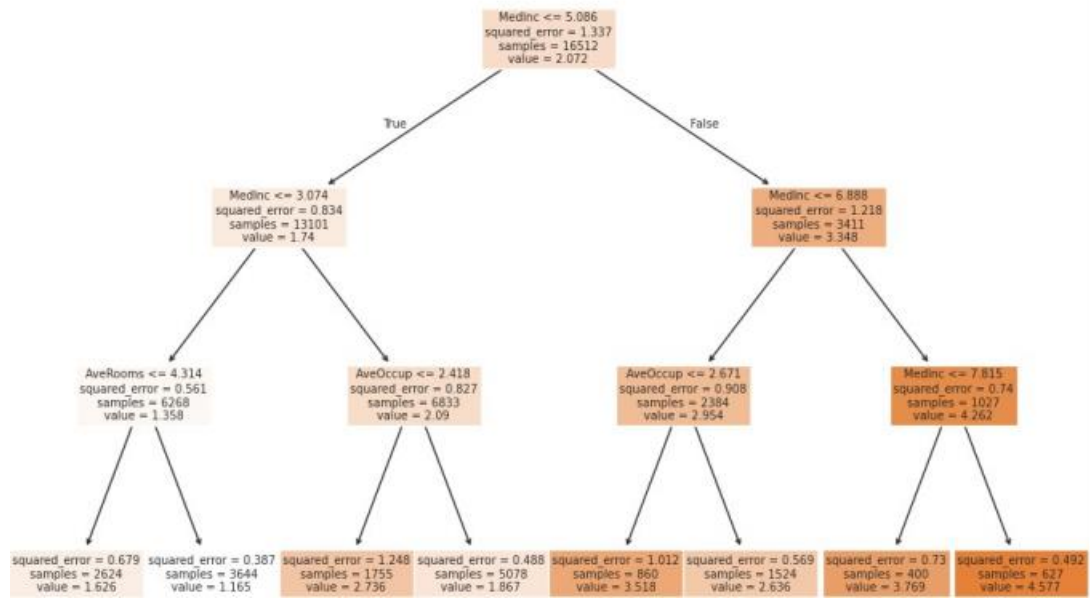
mse = mean_squared_error(y_test_housing, y_pred_housing)
print(f'Mean Squared Error: {mse:.2f}')

Mean Squared Error: 0.64

plt.figure(figsize=(12,8))
plot_tree(reg, filled=True, feature_names=housing.feature_names)
plt.title("Decision Tree Regressor")
plt.show()

```


Decision Tree Regressor



4.4 Implement a Support Vector Machine for any relevant dataset.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns

iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
svm_classifier = SVC(kernel='linear', random_state=42)
svm_classifier.fit(X_train, y_train)

y_pred = svm_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(conf_matrix)

class_report = classification_report(y_test, y_pred, target_names=iris.target_names)
print("\nClassification Report:")
print(class_report)
```

	19	0	0
	0	13	0
	0	0	13

Accuracy: 1.00

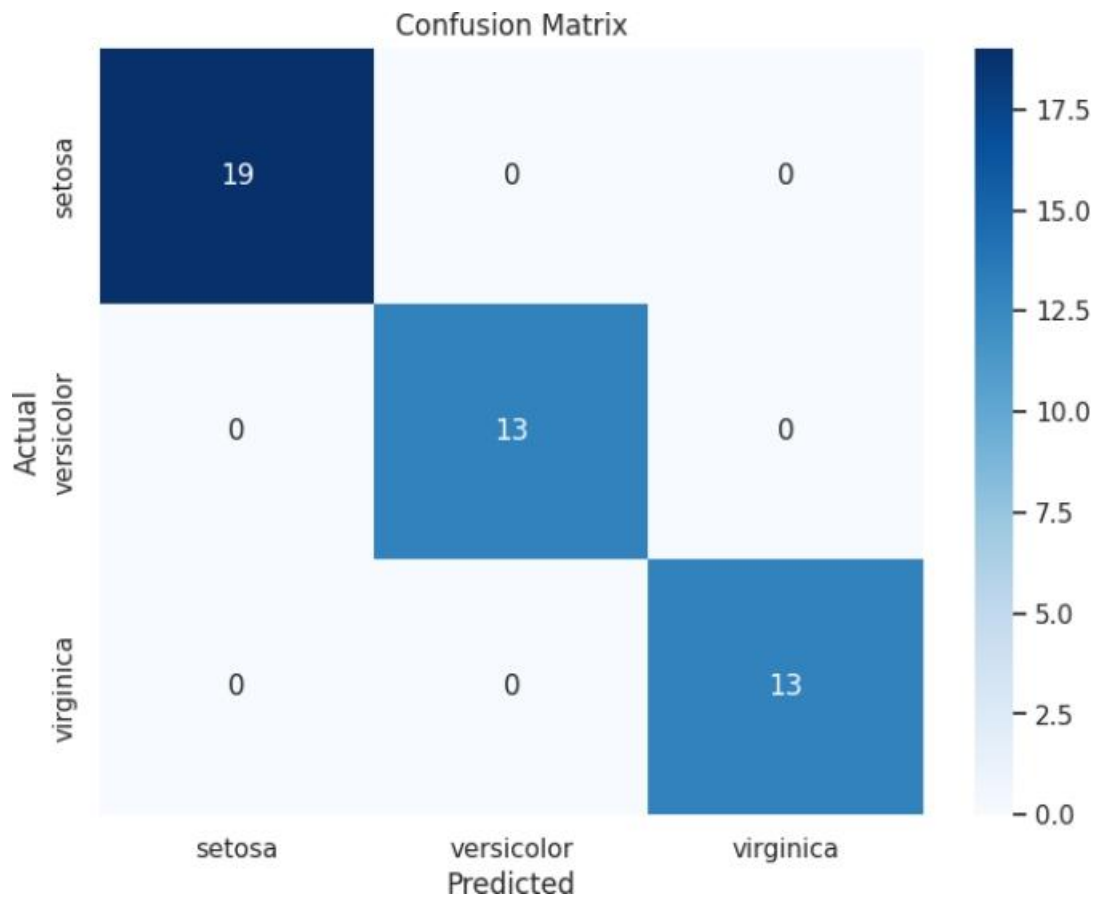
Confusion Matrix:

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13

```
accuracy          1.00          1.00          1.00          45
macro avg         1.00          1.00          1.00          45
weighted avg      1.00          1.00          1.00          45
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.title('Confusion Matrix')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



4.5 Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns

iris = datasets.load_iris()
X = iris.data
y = iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
tree_classifier = DecisionTreeClassifier(random_state=42)
tree_classifier.fit(X_train, y_train)
y_pred_tree = tree_classifier.predict(X_test)
accuracy_tree = accuracy_score(y_test, y_pred_tree)
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}\n")

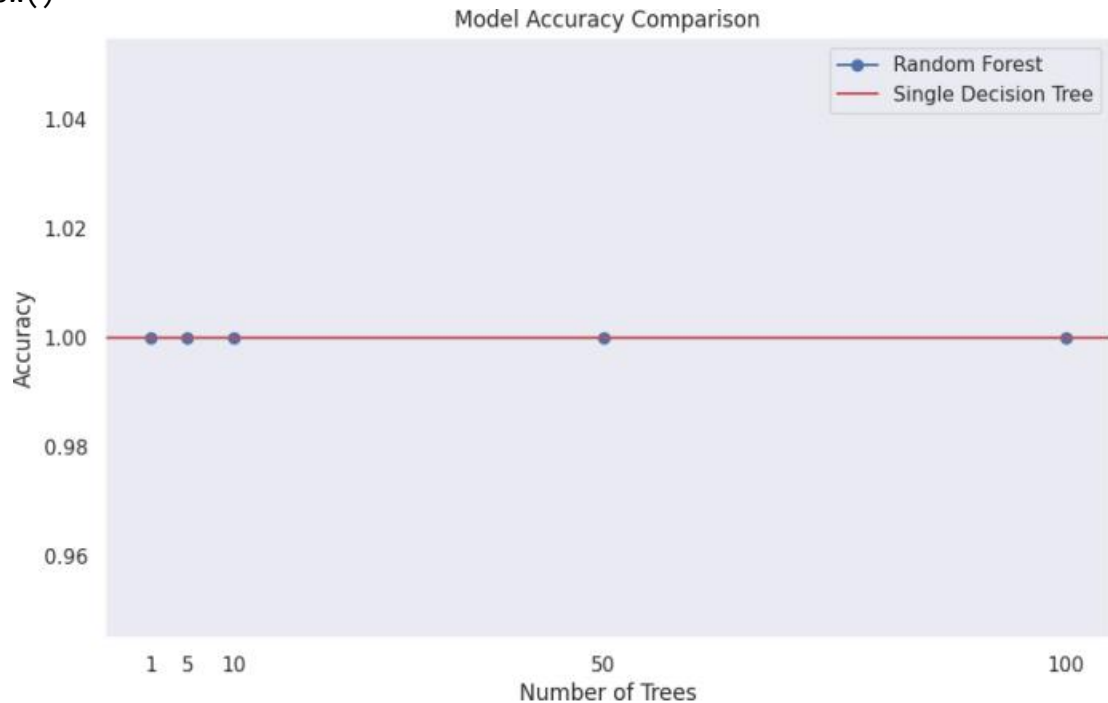
Decision Tree Accuracy: 1.00

n_trees = [1, 5, 10, 50, 100]
accuracy_forest = []
for n in n_trees:
    forest_classifier = RandomForestClassifier(n_estimators=n, random_state=42)
    forest_classifier.fit(X_train, y_train)
    y_pred_forest = forest_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred_forest)
    accuracy_forest.append(accuracy)
    print(f"Random Forest with {n} trees Accuracy: {accuracy:.2f}")

Random Forest with 1 trees Accuracy: 1.00
Random Forest with 5 trees Accuracy: 1.00
Random Forest with 10 trees Accuracy: 1.00
Random Forest with 50 trees Accuracy: 1.00
Random Forest with 100 trees Accuracy: 1.00

plt.figure(figsize=(10, 6))
plt.plot(n_trees, accuracy_forest, marker='o', label='Random Forest')
plt.axhline(y=accuracy_tree, color='r', linestyle='-', label='Single Decision Tree')
```

```
plt.title('Model Accuracy Comparison')
plt.xlabel('Number of Trees')
plt.ylabel('Accuracy')
plt.xticks(n_trees) plt.legend()
plt.grid()
plt.show()
```



```
best_n = n_trees[np.argmax(accuracy_forest)] # Get the best performing number of trees
best_forest_classifier = RandomForestClassifier(n_estimators=best_n, random_state=42)
best_forest_classifier.fit(X_train, y_train)
y_pred_best_forest = best_forest_classifier.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_best_forest)
print("\nConfusion Matrix for Random Forest (best model):")
print(conf_matrix)

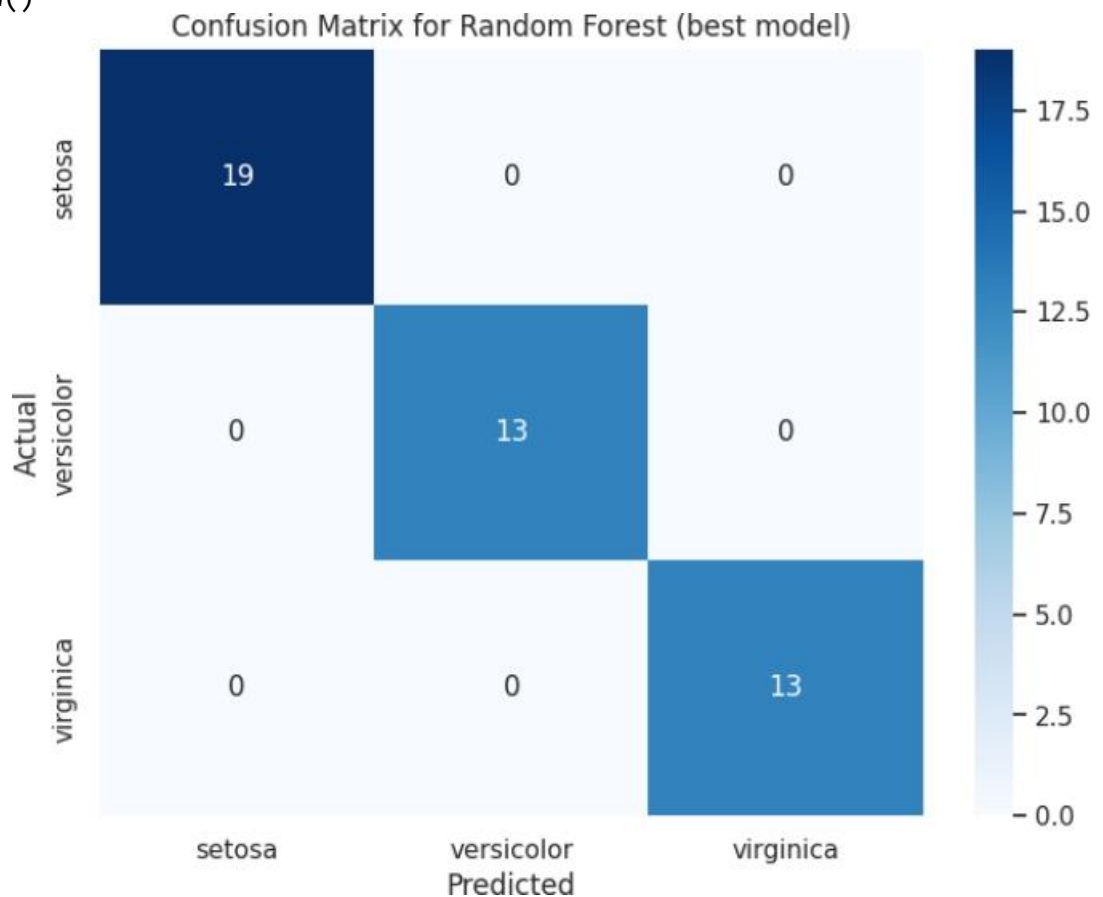
Class_report = classification_report(y_test, y_pred_best_forest, target_names=iris.target_names)
print("\nClassification Report for Random Forest (best model):")
print(class_report)

Confusion Matrix for Random Forest (best model):
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

Classification Report for Random Forest (best model):

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names, yticklabels=iris.target_names)
plt.title('Confusion Matrix for Random Forest (best model)')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



4.6 Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
from sklearn.model_selection import train_test_split, GridSearchCV
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb_model.fit(X_train, y_train)

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [3, 4, 5],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0]
}
grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, scoring='accuracy', cv=3, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
print("Best parameters from GridSearch:", best_params)

Fitting 3 folds for each of 54 candidates, totalling 162 fits
Best parameters from GridSearch: {'learning_rate': 0.01, 'max_depth': 4, 'n_estimators': 100, 'subsample': 0.8}

best_xgb_model = grid_search.best_estimator_
y_pred = best_xgb_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"\nXGBoost Accuracy: {accuracy:.2f}\n")

conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
```

```

class_report = classification_report(y_test, y_pred, target_names=iris.target_names)
print("\nClassification Report:")
print(class_report)

```

XGBoost Accuracy: 1.00

Confusion Matrix:

```

[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]

```


Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

```

plt.figure(figsize=(10, 6))
sns.barplot(x=best_xgb_model.feature_importances_, y=iris.feature_names)
plt.title('Feature Importance from XGBoost')
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.show()

```

Practical 5

5.1 Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample.

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
iris = load_iris()
X = iris.data
y = iris.target
```

```
df = pd.DataFrame(data=X, columns=iris.feature_names)
df['target'] = y
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.
2				
1	4.9	3.0	1.4	0.
2				
2	4.7	3.2	1.3	0.
2				
3	4.6	3.1	1.5	0.
2				
4	5.0	3.6	1.4	0.
2				


```
print(df.head())
```

```
sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
) \
0                5.1                3.5                1.4                0.
2
1                4.9                3.0                1.4                0.
2
2                4.7                3.2                1.3                0.
2
```

3	4.6	3.1	1.5	0.
2				
4	5.0	3.6	1.4	0.
2				

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, r
andom_state=42)
```

```
gnb      =      GaussianNB()
gnb.fit(X_train, y_train)
```

```
y_pred = gnb.predict(X_test)
```

```
accuracy      =      accuracy_score(y_test,      y_pred)
conf_matrix    =      confusion_matrix(y_test,      y_pred)
class_report   =      classification_report(y_test, y_pred)
```

```
print(f"Accuracy:      {accuracy:.2f}")
print("Confusion      Matrix:")
print(conf_matrix)
print("Classification  Report:")
print(class_report)
```

Accuracy: 0.98

Confusion Matrix:

```
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.92	0.96	13
2	0.93	1.00	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

5.2 Implement Hidden Markov Models using hmmlearn

```
!pip install hmmlearn
import numpy as np
import matplotlib.pyplot as plt
from hmmlearn import hmm

Requirement already satisfied: hmmlearn in /usr/local/lib/python3.10/dist-packages (0.3.2)
Requirement already satisfied: numpy>=1.10 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.26.4)
Requirement already satisfied: scikit-learn!=0.22.0,>=0.16 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.5.2)
Requirement already satisfied: scipy>=0.19 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn!=0.22.0,>=0.16->hmmlearn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn!=0.22.0,>=0.16->hmmlearn) (3.5.0)

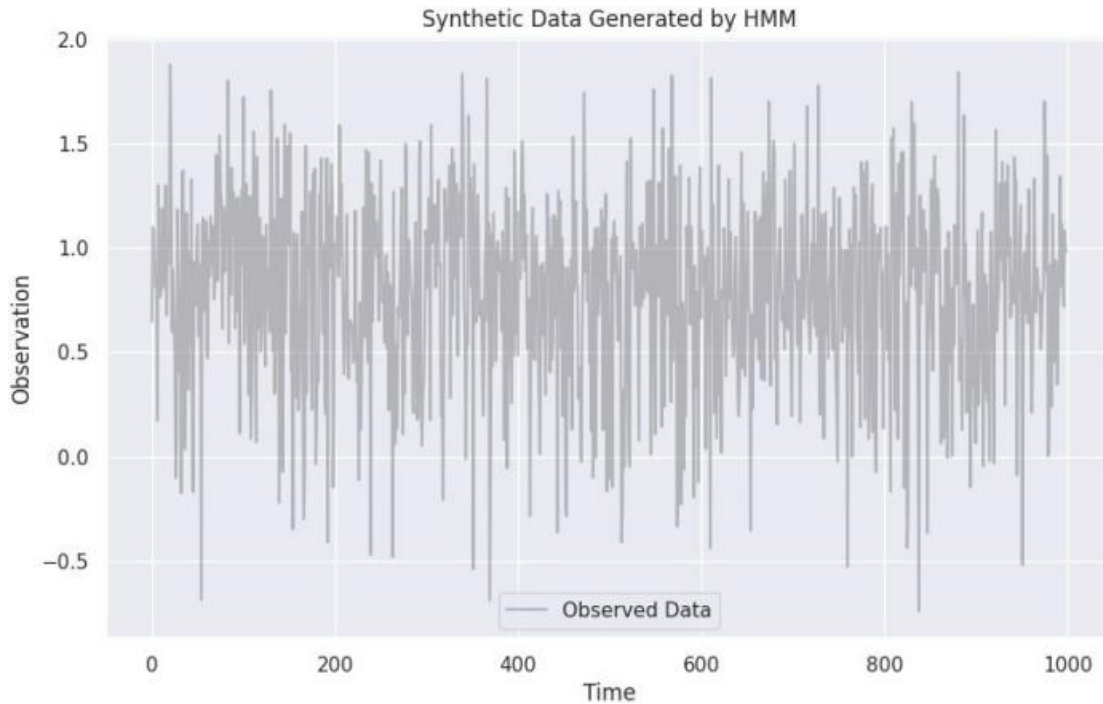
np.random.seed(42)

n_samples = 1000
n_states = 2
trans_probs = np.array([[0.7, 0.3],
                        [0.4, 0.6]])
means = np.array([[1.0], [0.5]])
covars = np.array([[0.1], [0.2]])

model = hmm.GaussianHMM(n_components=n_states, covariance_type="diag", n_iter=100)
model.startprob_ = np.array([0.6, 0.4])
model.transmat_ = trans_probs
model.means_ = means
model.covars_ = covars

X, Z = model.sample(n_samples)

plt.figure(figsize=(10, 6))
plt.plot(X, label='Observed Data', color='grey', alpha=0.5)
plt.title('Synthetic Data Generated by HMM')
plt.xlabel('Time')
plt.ylabel('Observation')
plt.legend()
plt.show()
```



```
model.fit(X)
hidden_states = model.predict(X)
```


WARNING:hmmlearn.base:Even though the 'startprob_' attribute is set, it will be overwritten during initialization because 'init_params' contains 's'

WARNING:hmmlearn.base:Even though the 'transmat_' attribute is set, it will be overwritten during initialization because 'init_params' contains 't'

WARNING:hmmlearn.base:Even though the 'means_' attribute is set, it will be overwritten during initialization because 'init_params' contains 'm'

WARNING:hmmlearn.base:Even though the 'covars_' attribute is set, it will be overwritten during initialization because 'init_params' contains 'c'

```
plt.figure(figsize=(10, 6))
plt.plot(X, label='Observed Data', color='grey', alpha=0.5)
plt.step(range(n_samples), hidden_states, where="post", label='Hidden States', color='red')
plt.title('Hidden States Predicted by HMM')
plt.xlabel('Time')
plt.ylabel('State')
plt.legend()
plt.show()
```



```
print("Transition matrix:\n", model.transmat_)
print("Means:\n", model.means_)
print("Covariances:\n", model.covars_)
```

```
Transition      matrix:
[[0.65865532 0.34134468]
 [0.3121865  0.6878135  ]]
```

```
Means:
[[0.54954006]
 [1.00338912]]
```

```
Covariances:
[[[0.22176075]]

 [[0.09283459]]]
```

Practical 6


6.1 Implement Bayesian Linear Regression to explore prior and posterior distribution.

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal

np.random.seed(42)
X = np.random.rand(100, 1) * 10
true_beta = np.array([2.0])
y = 2.0 * X.flatten() + np.random.normal(0, 2, size=X.shape[0])
X_b = np.c_[np.ones((X.shape[0], 1)), X]
sigma_0 = 10
sigma_n = 4
sigma_0_inv = 1 / sigma_0
sigma_n_inv = 1 / sigma_n
N = X_b.shape[0]
beta_prior_mean = np.zeros(X_b.shape[1])
beta_prior_cov = sigma_0 * np.eye(X_b.shape[1])
posterior_cov = np.linalg.inv(sigma_n_inv * (X_b.T @ X_b) + sigma_0_inv *
np.eye(X_b.shape[1]))
posterior_mean = posterior_cov @ (sigma_n_inv * (X_b.T @ y))
beta_samples = np.random.multivariate_normal(posterior_mean, posterior_cov
, size=1000)
plt.figure(figsize=(10, 6))
beta_prior_samples = np.random.multivariate_normal(beta_prior_mean, beta_p
rior_cov, size=1000)

<Figure size 1000x600 with 0 Axes>

plt.subplot(1, 2, 1)
plt.title("Prior Distribution")
plt.scatter(beta_prior_samples[:, 0], beta_prior_samples[:, 1], alpha=0.2)
plt.xlabel("$\\beta_0$")
plt.ylabel("$\\beta_1$")
plt.subplot(1, 2, 2)
plt.title("Posterior Distribution")
plt.scatter(beta_samples[:, 0], beta_samples[:, 1], alpha=0.2, color='oran
ge')
plt.xlabel("$\\beta_0$")
plt.ylabel("$\\beta_1$")
plt.tight_layout()
plt.show()
```



```
print("Posterior      Mean:",      posterior_mean)
print("Posterior Covariance:\n", posterior_cov)
```


```
Posterior      Mean:      [0.42825291      1.90809351]
Posterior Covariance:
[[ 0.1389247 -0.0211579 ]
 [-0.0211579  0.00451795]]
```


6.2 Implement Gaussian Mixture Models for density estimation and unsupervised clustering

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture

np.random.seed(42)
means = [[2, 2], [8, 8], [5, 1]]
covariances = [[[1, 0], [0, 1]], [[1, 0], [0, 1]], [[1, 0], [0, 1]]]
n_samples = 500
data = np.vstack([
    np.random.multivariate_normal(mean, cov, n_samples // len(means))
    for mean, cov in zip(means, covariances)
])
n_components = len(means) # Number of clusters
gmm = GaussianMixture(n_components=n_components, covariance_type='full')
gmm.fit(data)
labels = gmm.predict(data)

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.scatter(data[:, 0], data[:, 1], c=labels, s=30, cmap='viridis', alpha=0.5)
plt.title('GMM Clustering')
plt.xlabel('X1')
plt.ylabel('X2')
x = np.linspace(-1, 10, 100)
y = np.linspace(-1, 10, 100)
X, Y = np.meshgrid(x, y)
XX = np.column_stack([X.ravel(), Y.ravel()])
logprob = gmm.score_samples(XX)
pdf = np.exp(logprob).reshape(X.shape)
plt.subplot(1, 2, 2)
plt.contourf(X, Y, pdf, levels=20, cmap='Blues', alpha=0.7)
plt.scatter(data[:, 0], data[:, 1], c='red', s=10, alpha=0.5)
plt.title('Density Estimation using GMM')
plt.xlabel('X1')
plt.ylabel('X2')
plt.colorbar(label='Density')
plt.tight_layout()
plt.show()
```



Practical 7

7.1 Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation

```
import numpy as np
from sklearn.model_selection import KFold, StratifiedKFold
from sklearn.metrics import accuracy_score
from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier

data = load_iris()
X, y = data.data, data.target
model = RandomForestClassifier()
kf = KFold(n_splits=5, shuffle=True, random_state=42)
kf_scores = []
for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    score = accuracy_score(y_test, predictions)
    kf_scores.append(score)
print(f'K-Fold Accuracy: {np.mean(kf_scores):.2f} ± {np.std(kf_scores):.2f}')

K-Fold Accuracy: 0.96 ± 0.02

skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
skf_scores = []
for train_index, test_index in skf.split(X, y):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    score = accuracy_score(y_test, predictions)
    skf_scores.append(score)
print(f'Stratified K-Fold Accuracy: {np.mean(skf_scores):.2f} ± {np.std(skf_scores):.2f}')

Stratified K-Fold Accuracy: 0.95 ± 0.03
```

7.2 Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split, GridSearchCV, Random
izedSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

data = load_iris()
X, y = data.data, data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=42)
model = RandomForestClassifier(random_state=42)
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
}
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, scoring
='accuracy', cv=5)
grid_search.fit(X_train, y_train)
print("Grid Search Best Parameters:", grid_search.best_params_)
print("Grid Search Best Score:", grid_search.best_score_)

Grid Search Best Parameters: {'max_depth': 5, 'min_samples_split': 5, 'n_e
stimators': 10}
Grid Search Best Score: 0.9636363636363636

param_dist = {
    'n_estimators': np.arange(10, 200, 10),
    'max_depth': [None] + list(np.arange(1, 20, 1)),
    'min_samples_split': np.arange(2, 20, 2),
}
random_search = RandomizedSearchCV(estimator=model, param_distributions=pa
ram_dist, n_iter=50, scoring='accuracy', cv=5, random_state=42)
random_search.fit(X_train, y_train)
print("Randomized Search Best Parameters:", random_search.best_params_)
print("Randomized Search Best Score:", random_search.best_score_)

Randomized Search Best Parameters: {'n_estimators': 120, 'min_samples_spli
t': 16, 'max_depth': None}
Randomized Search Best Score: 0.9636363636363636

best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
```

```
test_accuracy = accuracy_score(y_test, y_pred)
print("Test Set Accuracy using Grid Search Best Model:", test_accuracy)
```

```
Test Set Accuracy using Grid Search Best Model: 1.0
```

Practical 8

8.1 Implement Bayesian Learning using inferences

```
import numpy as np

P_A = 0.5
P_B = 0.5

def likelihood_heads(coin, flips):
    if coin == 'A':
        return (0.5 * flips) * (0.5 * (10 - flips))
    elif coin == 'B':
        return (0.9 * flips) * (0.1 * (10 - flips))

observed_heads = 8
total_flips = 10

likelihood_A = likelihood_heads('A', observed_heads)
likelihood_B = likelihood_heads('B', observed_heads)

marginal_likelihood = (likelihood_A * P_A) + (likelihood_B * P_B)

posterior_A = (likelihood_A * P_A) / marginal_likelihood
posterior_B = (likelihood_B * P_B) / marginal_likelihood

print(f"Posterior Probability of Coin A: {posterior_A:.4f}")
print(f"Posterior Probability of Coin B: {posterior_B:.4f}")

Posterior Probability of Coin A: 0.7353
Posterior Probability of Coin B: 0.2647
```

Practical 9

9.1 Set up a generator network to produce samples and a discriminator network to distinguish between real and generated data. (Use a simple small dataset)

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers, models

(X_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
X_train = (X_train.astype(np.float32) - 127.5) / 127.5
X_train = np.expand_dims(X_train, axis=-1)

latent_dim = 100
num_examples_to_generate = 16

def build_generator():
    model = models.Sequential()
    model.add(layers.Dense(256, input_dim=latent_dim))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(512))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(1024))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(28 * 28 * 1, activation='tanh'))
    model.add(layers.Reshape((28, 28, 1)))
    return model

def build_discriminator():
    model = models.Sequential()
    model.add(layers.Flatten(input_shape=(28, 28, 1)))
    model.add(layers.Dense(512))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Dense(256))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Dense(1, activation='sigmoid'))
    return model

generator = build_generator()
discriminator = build_discriminator()
discriminator.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```

discriminator.trainable = False
gan_input = layers.Input(shape=(latent_dim,))
generated_image = generator(gan_input)
gan_output = discriminator(generated_image)
gan = models.Model(gan_input, gan_output)
gan.compile(optimizer='adam', loss='binary_crossentropy')

def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input)
    predictions = (predictions.numpy() + 1) / 2 # Rescale to [0, 1]

    plt.figure(figsize=(4, 4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i + 1)
        plt.imshow(predictions[i, :, :, 0], cmap='gray')
        plt.axis('off')
    plt.savefig(f'gan_epoch_{epoch}.png')
    plt.show()

def train_gan(epochs, batch_size):
    random_latent_vectors = tf.random.normal(shape=(num_examples_to_generate, latent_dim))

    for epoch in range(epochs):
        idx = np.random.randint(0, X_train.shape[0], batch_size)
        real_images = X_train[idx]

        noise = tf.random.normal(shape=(batch_size, latent_dim))
        fake_images = generator(noise)

        combined_images = tf.concat([real_images, fake_images], axis=0)

        labels = tf.constant([[1.0]] * batch_size + [[0.0]] * batch_size)

        d_loss = discriminator.train_on_batch(combined_images, labels)

        noise = tf.random.normal(shape=(batch_size, latent_dim))
        misleading_labels = tf.constant([[1.0]] * batch_size)

        g_loss = gan.train_on_batch(noise, misleading_labels)

        if epoch % 100 == 0:
            print(f"Epoch: {epoch}")
            print(f"Discriminator Loss: {d_loss[0]}")
            print(f"Generator Loss: {g_loss}")

        generate_and_save_images(generator, epoch, random_latent_vectors)

```




```
epochs = 300  
batch_size = 64
```

```
train_gan(epochs, batch_size)
```

Epoch: 0

Discriminator Loss: 0.7258248329162598


Generator Loss: [array(0.72582483, dtype=float32), array(0.72582483, dtype=float32), array(0.390625, dtype=float32)]



Epoch: 100

Discriminator Loss: 2.1150412559509277


Generator Loss: [array(2.1150413, dtype=float32), array(2.1150413, dtype=float32), array(0.20482673, dtype=float32)]



Epoch: 200

Discriminator Loss: 2.8626105785369873

Generator Loss: [array(2.8626106, dtype=float32), array(2.8626106, dtype=float32), array(0.20747824, dtype=float32)]



Practical 10

10.1 Develop an API to deploy your model and perform predictions

```
# Required Libraries
!pip install pyngrok flask scikit-learn

# Importing Libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import pickle
from flask import Flask, request, jsonify
from pyngrok import ngrok

# Load dataset
iris = load_iris()
X, y = iris.data, iris.target

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
andom_state=42)

# Train a model
model = RandomForestClassifier()
model.fit(X_train, y_train)

# Save the model
with open('model.pkl', 'wb') as model_file:
    pickle.dump(model, model_file)

# Load the model
with open('model.pkl', 'rb') as model_file:
    model = pickle.load(model_file)

# Create Flask app
app = Flask(__name__)
port = "5000"

@app.route('/')
def home():
    return "Welcome to the Iris Prediction API! Use the /predict endpoint
to make predictions."

@app.route('/predict', methods=['POST'])
def predict():
    data = request.json
    features = data.get('features')
```

```

    # Ensure the features are in the correct format
    if not features or len(features) != 4: # Assuming 4 features for
iris dataset
        return jsonify({'error': 'Invalid input format. Please
provide 4 features.'}), 400

    try:
        prediction = model.predict([features]) # Wrap features in a
        list
        to create 2D array
        return jsonify({'prediction': int(prediction[0])}) # Convert
        prediction to int
    except Exception as e:
        return jsonify({'error': str(e)}), 500

# Start ngrok and print the public URL
ngrok.set_auth_token("api_auth_token")
public_url =
ngrok.connect(port).public_url
print("Public URL:", public_url)

# Run the Flask app
if __name__ == '__main__':
    app.run(port=port)

```

```

Requirement          already          satisfied:          pyngrok          in
/usr/local/lib/python3.10/dist-packages (7.2.0)
Requirement          already          satisfied:          flask          in
/usr/local/lib/python3.10/dist-packages (2.2.5)
Requirement          already          satisfied:          scikit-learn          in
/usr/local/lib/python3.10/dist-packages (1.5.2)
Requirement          already          satisfied:          PyYAML>=5.1          in
/usr/local/lib/python3.10/dist-packages (from pyngrok) (6.0.2)
Requirement          already          satisfied:          Werkzeug>=2.2.2          in
/usr/local/lib/python3.10/dist-packages (from flask) (3.0.4)
Requirement          already          satisfied:          Jinja2>=3.0          in
/usr/local/lib/python3.10/dist-packages (from flask) (3.1.4)
Requirement          already          satisfied:          itsdangerous>=2.0          in
/usr/local/lib/python3
.10/dist-packages (from flask) (2.2.0)
Requirement          already          satisfied:          click>=8.0          in
/usr/local/lib/python3.10/dist-packages (from flask) (8.1.7)
Requirement          already          satisfied:          numpy>=1.19.5          in
/usr/local/lib/python3.10/
dist-packages (from scikit-learn)
(1.26.4)
Requirement          already          satisfied:          scipy>=1.6.0          in
/usr/local/lib/python3.10/d
ist-packages (from scikit-learn)
(1.13.1)
Requirement          already          satisfied:          joblib>=1.2.0          in
/usr/local/lib/python3.10/
dist-packages (from scikit-learn) (1.4.2)
Requirement          already          satisfied:          threadpoolctl>=3.1.0          in
/usr/local/lib/pyth on3.10/dist-packages (from scikit-learn) (3.5.0)
Requirement          already          satisfied:          MarkupSafe>=2.0          in

```

```
/usr/local/lib/python3.10/dist-packages (from Jinja2>=3.0->flask)
(2.1.5)
Public URL: https://2f62-49-43-24-101.ngrok-free.app
```

```
* Serving Flask app '__main__'
* Debug mode: off
```

```
INFO:werkzeug:WARNING: This is a development server. Do not use it in a
production deployment. Use a production WSGI server instead.
```

```
* Running on
```

```
http://127.0.0.1:5000
```

```
INFO:werkzeug:Press CTRL+C to
```

```
quit
```

```
INFO:werkzeug:127.0.0.1 - - [24/Oct/2024 07:50:42] "POST /predict
HTTP
```